

# 60017 PERFORMANCE ENGINEERING

## Bottlenecks in Distributed Systems

# Last lecture

- ▶ Workloads in computer systems
  - ▶ Log files
  - ▶ User behavior models


# This lecture

- ▶ Service demand
- ▶ Utilization law
- ▶ Bottleneck analysis

# Overview

- ▶ Typical questions involving computational resources include:
  - ▶ **Bottleneck analysis**: Which resources limit the scalability of a given distributed application?
  - ▶ **What-if analysis**: How will changes in user request rates or resource speed impact performance?
- ▶ Performance of individual resources (e.g., CPUs) affected by:
  - ▶ **Arrival rate** of requests
  - ▶ **Service time** of the request at the resource
  - ▶ **Contention** at the resource (e.g., jobs contend the CPU)
  - ▶ ...
- ▶ How to predict and analyse resource utilizations?

## Example: resources in a three-tier application



Source: Wang et al, CLOUD 2013

- ▶ Serving a web page request typically involves multiple servers
- ▶ Each server (e.g., MySQL) is a resource
- ▶ A request can visit multiple times a resource, receiving a variable **service time** at each visit, until completing.
- ▶ We here focus on the percentage of time each resource is **busy** and do not attempt to quantify overheads due to **contention**.

# Service classes

- ▶ We assume that the IT system:
  - ▶ offers  $C$  types of services (**service classes**), e.g.,  $C$  web pages.
  - ▶ processes requests using  $M$  resources, e.g.,  $M$  servers.
- ▶ In order to complete, a request of class  $c$ 
  - ▶ makes on average  $k_{ic}$  calls to resource  $i$
  - ▶ requires a mean service time of  $S_{ic}$  seconds for each call to  $i$
- ▶ The **demand**  $D_{ic} = k_{ic}S_{ic}$  is the service time accumulated by a class- $c$  request through its visits to resource  $i$ .
- ▶ In the presence of contention/queueing, the demand represents only the effective processing time received at the resource, neglecting overheads due to other jobs.

# Operational analysis

- ▶ Let us now investigate the relationship between demands and resource utilization.
- ▶ Suppose to monitor a resource for an **observation period** of  $T$  seconds collecting:
  - ▶  $A_c$ : total number of **arrived** requests of class  $c$
  - ▶  $B_{ic}$ : total time resource  $i$  is **busy** processing class- $c$  requests
- ▶ Averaging over the observation period we obtain:
  - ▶  $\lambda_c = A_c/T$ : average **arrival rate** of requests of class  $c$
  - ▶  $U_{ic} = B_{ic}/T$ : **utilization** of resource  $i$  due to class- $c$  requests

# Operational analysis

- ▶ If requests never fail and the pending requests are always bounded in number, then  $\lambda_c$  is also equal to the class- $c$  **throughput**  $X_c$ , as  $T$  grows large because

$$\lambda_c = \lim_{T \rightarrow \infty} \frac{A_c}{T} = \lim_{T \rightarrow \infty} \frac{\text{completed}(c) + \text{pending}(c)}{T} = \lim_{T \rightarrow \infty} \frac{\text{completed}(c)}{T} = X_c$$


since the number of pending requests remains finite, if the system is **stable**.

- ▶ An **unstable** system cannot cope with the rate of arrivals, so that eventually its backlog grows unbounded.
- ▶ Throughout, we assume that the system is stable.



## Example: busy time, arrivals, completions

Assume a single workload class ( $C = 1$ ). We omit class indices.



$$B_{i1} = 7 + 7 + 2 + 10 + 5 + 3 + 6 = 40s$$

$$U_i = B_{i1}/T = 40/58 = 0.689 \text{ (68.9\%)}$$

$$\lambda = \text{upward transitions}/T = 5/58 \text{ tps}$$

$$X = \text{downward transitions}/T = 5/58 \text{ tps} = \lambda$$

# Utilization law

- ▶ We now observe that

$$D_{ic} = B_{ic}/A_c$$

since  $B_{ic}/A_c$  averages the total processing time of class- $c$  requests at  $i$  (which is the definition of  $D_{ic}$ ).

- ▶ What is the relationship between utilization and demands?


$$U_{ic} = \frac{B_{ic}}{T} = \frac{A_c}{T} \frac{B_{ic}}{A_c} = \lambda_c D_{ic}$$

- ▶ For a system with  $C$  workload classes, the total resource utilization  $U_i$  is thus given by the **utilization law**

$$U_i = \sum_{c=1}^C U_{ic} = \sum_{c=1}^C \lambda_c D_{ic} = \sum_{c=1}^C X_c D_{ic}$$

## Example: utilization from UBGs

- Users arrive at rate  $\lambda = 1$  sessions/sec



$D_{ic}$ [sec]	Home	Search	Add	Buy
Web	0.1	0.2	0.4	0.1
Tomcat	0.0	0.2	0.3	0.7
MySQL	0.0	0.8	0.2	0.1

- UBG visit ratios:

$$V_H = 1, V_A = 0.725, V_S = 1.75, V_B = 0.29$$

- Request rates:

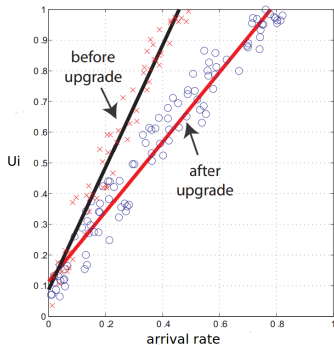
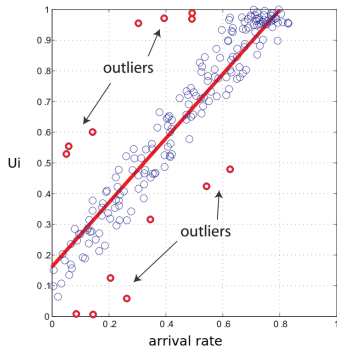
$$\lambda_H = \lambda V_H, \lambda_A = \lambda V_A, \lambda_S = \lambda V_S, \lambda_B = \lambda V_B$$

- Utilization:

$$U_W = \lambda_H D_{WH} + \lambda_A D_{WA} + \lambda_S D_{WS} + \lambda_B D_{WB} = 0.974 \text{ (97.4\%)}$$

# Demand estimation from utilization samples

- ▶ How to estimate demand values in practice?
- ▶ The utilization law describes a hyperplane with slopes  $D_{ic}$ .



- ▶ **Multivariate linear regression** fits hyperplanes to samples of  $U_i$  and  $\lambda_c$  (or  $X_c$ ),  $\forall c$ , returning the demands  $D_{ic}$ .
- ▶ The estimated demand values are hardware dependent and in general change after hardware upgrades.

# Bottleneck analysis

- ▶ Demands can be used to determine the resources that limit scalability, called the **bottlenecks**.
- ▶ The bottlenecks are **oversubscribed** resources, which struggle to complete the backlog of pending requests.
- ▶ Bottlenecks often run near 100% utilization (**saturation**), thus over time they tend to pile up a large backlog of requests.
- ▶ **Bottleneck analysis** wishes to answer these questions:
  - ▶ As we increase the arrival rates of requests, which resource(s) will **saturate** first?
  - ▶ What is the **maximum arrival rate** that the system can sustain under a given request mix?

## Bottleneck analysis

- ▶ For a system with  $M$  resources, we can describe resource usage in terms of the system of linear equations

$$U_i = \sum_{c=1}^C \lambda_c D_{ic} \quad i = 1, \dots, M$$

where we require  $U_i \leq 1$  for each resource.


- ▶ Resource  $j$  can saturate if and only if there exists a combination of arrival rates  $(\lambda_1, \dots, \lambda_C)$  such that  $U_j = 1$ .
- ▶ Whether a resource can saturate may be verified using the following **linear program (LP)** defined over the variables  $\lambda_c$

$$\begin{aligned} U_j^{\max} = \text{maximize} \quad & \sum_{c=1}^C \lambda_c D_{jc} \\ \text{subject to} \quad & \sum_{c=1}^C \lambda_c D_{ic} \leq 1, \quad i = 1, \dots, M \\ & \lambda_c \geq 0, \quad c = 1, \dots, C \end{aligned}$$

If the LP has optimal value  $U_j^{\max} = 1$ , then  $j$  can saturate.

## Example: sustainable arrival rates

Assume  $C = 2$  workload classes and  $M = 4$  resources.



Each line represents the boundary of  $U_i \leq 1$ , for given  $i$  and  $D_{ic}$ .

## Example: sustainable arrival rates

In the example:

- ▶ Each point  $(\lambda_1, \lambda_2)$  is a possible arrival rate to the system.
- ▶ The shaded region indicates the **sustainable arrival rates**.
- ▶ Server 1 can never become a bottleneck. Never upgrade it!
- ▶ Server 2 is the **class-1 bottleneck** (largest demand in class 1).
- ▶ Server 3 is the **class-2 bottleneck** (largest demand in class 2).
- ▶ Server 4 is a bottleneck only for some arrival rates.
- ▶ Servers 2,3,4 are thus **potential bottlenecks** for this system.
- ▶ Some mixes lead multiple resources to become bottlenecks **simultaneously**, e.g., 2 + 3 or 2 + 4.





# Bottleneck mitigation

- ▶ If the system has a single potential bottleneck, we can **remove the bottleneck** by upgrading its hardware or scale up the VM.
- ▶ However, it is important to realise that arrival rates **change over time** and so the active bottlenecks.
  - ▶ Arrival rates change visibly through the day, week and month.
  - ▶ Recent studies found that the class mix also varies frequently.
- ▶ Therefore IT systems may display different bottlenecks over time, requiring **adaptive resource management** to cope with transient need of capacity at a certain resource.
  - ▶ Since cloud computing allows to acquire and release resources via APIs, it facilitates adaptive resource management.

## Example: time-varying transaction mixes

Mix of requests in execution in a real e-commerce website:




Source: Stewart et al, Eurosys 2007

- ▶ Each point represent the fraction of requests of the two popular transaction types seen in the website during a 5-minute monitoring interval
- ▶ If the two requests stress different resources (e.g., Tomcat, MySQL), the resource usage of these resources will also fluctuate over time

# Bottleneck switches


- ▶ If there are multiple potential bottlenecks, rapid changes in the transaction mix can lead to observe **bottleneck switches**.
- ▶ Bottlenecks can migrate between resources very rapidly, even within milliseconds (**millibottlenecks**).
- ▶ If bottleneck switches occur at *ms* scale, they can be **invisible to monitoring**, which aggregates data at *s* or *min* resolution.



Sources: Wang et al., CLOUD 2013

# Diagnosing bottleneck switches

- ▶ Rapid bottleneck switches typically **degrade performance**.
  - ▶ Transactions more likely to find queues along the flow.
- ▶ A **symptom** that this is happening is that throughput grows slowly, even if the system is lightly utilised.
- ▶ Example: TPC-W's browsing mix illustrates rapid bottleneck migration between application server and DB.



Sources: Casale et al., IEEE TSE 2012