

60017 PERFORMANCE ENGINEERING

User Behaviour Modelling

Last lecture

- ▶ Benchmarking a distributed application
 - ▶ Case study: SPECjbb2015
- ▶ Load testing a distributed application

This lecture

- ▶ Workloads in computer systems
 - ▶ Log files
 - ▶ User behavior models


Log files and traces

- ▶ Workloads are recorded in **log files**. Log files record events at operating system or application level and their timestamps.
- ▶ Many tools automate log collection, filtering and analysis.
 - ▶ *e.g., Elasticsearch, Logstash, Kibana, ...*
- ▶ In distributed applications multiple log files exist and need to be combined with appropriate filters into **workload traces**.
 - ▶ *e.g., HTTP log files, log4j files, Linux /var/log files,*
- ▶ A workload trace organizes events that pertain to arrival and service of requests in a **time series**.

Log files and traces

- ▶ Traces can be used by managers and administrators to:
 - ▶ **understand** how the system is used by customers
 - ▶ **simulate** the system to answer what-if questions
 - ▶ **replay** a sequence of requests using a load testing tool

Example: arrival times of requests at a network-attached storage



Example: HTTP log files in Apache web servers

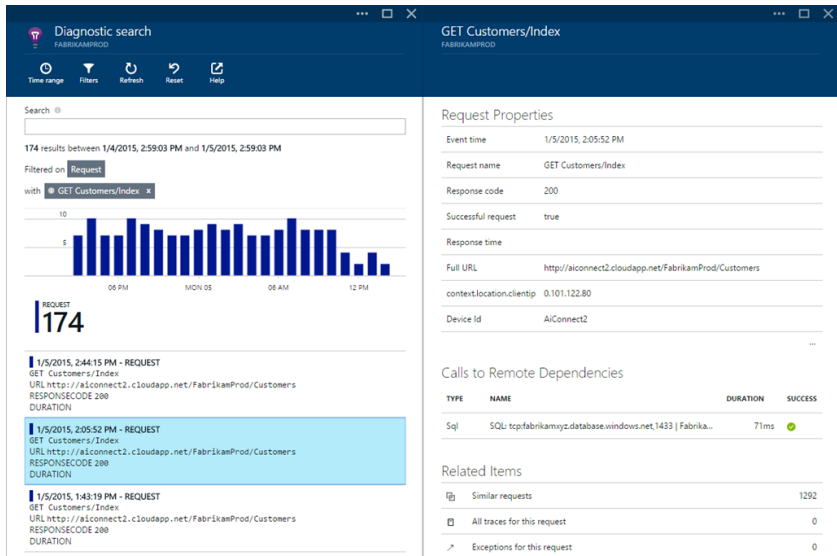
A typical Apache web server access log:

```
125.45.000.166 - - [10/Jun/2010:11:02:34 +0000] "GET http://www.yyy.it/index.html HTTP/1.0" 200 8859
"http://www.google.com/" "Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.1)"
125.45.000.166 - - [10/Jun/2010:11:04:17 +0000] "GET http://www.yyy.it/contact.html HTTP/1.1" 404
1010 "http://www.yyy.it/index.html" "Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.1)"
125.45.000.166 - - [19/Jun/2010:11:04:18 +0000] "GET http://www.yyy.it/contact.gif HTTP/1.1" 200 15890
 "-" "Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.1)"
...
```

Anatomy of the logfile format:

- ▶ 125.45.000.166: **client IP** address of the request
- ▶ [10/Jun/2010:11:02:34 +0000]: **timestamp**
- ▶ "GET http://www.yyy.it/index.html HTTP/1.0": **HTTP request** (GET for downloads, POST for forms)
- ▶ 200: **success code** (404 or 500 for errors)
- ▶ 8859: **size of response** in bytes, excluding HTTP headers.
- ▶ "http://www.google.com/": **referrer URL** (last page visited)
- ▶ "Mozilla/4.0 ...": **browser information**

Example: Azure AppInsights



Why modelling computer workloads?


- ▶ Workload traces are **essential** to understand IT system performance, but some **common issues** limit their use:
 - ▶ **Privacy**: companies don't like to give access (even to their own employees) to logs tracing customer activities.
 - ▶ **Inflexibility**: difficult to manipulate the properties of a trace to study system sensitivity to a workload parameter.
 - ▶ **Noise**: in a trace, we do not always know what is noise and what is not. This can mislead the analysis.
 - ▶ **Overfitting**: what we learn may depend too much on specific trace instances.

Why modelling computer workloads?

- ▶ **Workload model**: a model that can generate traces similar to the ones observed in the system.
 - ▶ *e.g.: statistical distributions, Markov chains, automata, ...*
- ▶ **Workload characterization**: model parameters are **fitted** to traces to capture their essential characteristics.
- ▶ Workload models have some advantages compared to traces:
 - ▶ **Repetitions**: models can generate similar, but nevertheless non-identical, workload traces.
 - ▶ **Understanding**: modelling increases our understanding and can lead to system optimizations based on it.
 - ▶ **Availability**: traces are not always available, but they can be generated by models instantiated with artificial parameters.
- ▶ We study some models for computer system workloads.

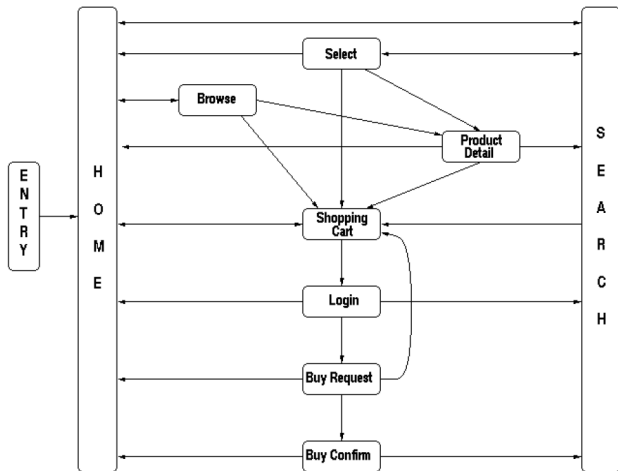
User behaviour graph (UBG)

- ▶ **User session**: the sequence of pages visited by a user.
- ▶ **UBG**: a probabilistic automaton describing user sessions.
 - ▶ **states** = page, session or service invocations
 - ▶ **arc weights** = transition probabilities
 - ▶ **outgoing arcs** = end of the session
- ▶ An example of UBG:



Example: UBG in TPC-W benchmark

- ▶ TPC-W: A classic benchmark for web servers (now retired).



User behaviour graph (UBG)

- ▶ UBGs are special **discrete-time Markov chains (DTMCs)**.
 - ▶ Users always start from the **Entry** state.
 - ▶ p_{ij} is the transition probability from node i to j
 - ▶ After visiting i , the user visits page j with probability p_{ij} .
 - ▶ Upon closing the session the user reaches the **Exit** node (not shown in the UBG).
- ▶ We denote by \mathcal{S} the set of DTMC states.

User behaviour graph (UBG)

- ▶ Every DTMC is described by a **transition probability matrix**

$$P = [p_{ij}] = \begin{matrix} & \begin{matrix} E & H & S & A & B & X \end{matrix} \\ \begin{matrix} E \\ H \\ S \\ A \\ B \\ X \end{matrix} & \left(\begin{array}{cccccc} 0 & 1.0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0.7 & 0.2 & 0 & 0.1 \\ 0 & 0 & 0.6 & 0.3 & 0 & 0.1 \\ 0 & 0 & 0 & 0 & 0.4 & 0.6 \\ 0 & 0 & 0 & 0 & 0 & 1.0 \\ 0 & 0 & 0 & 0 & 0 & 1.0 \end{array} \right) \end{matrix}$$

- ▶ Note: rows must sum to 1.
- ▶ Upon reaching the **Exit** (X) we never jump to another state. In DTMC theory this is called an **absorbing state**.

Using UBGs in practice

- ▶ UBGs can be fitted on data of a single or multiple client IPs
- ▶ p_{ij} values can be initially extracted from log files using the HTTP request and referer URL fields, e.g.,

$$p_{ij} = \frac{\text{number of requests for } j \text{ with } i \text{ as referer URL}}{\sum_{k \in \mathcal{S}} \text{number of requests for } k \text{ with } i \text{ as referer URL}}$$

- ▶ Care should be taken when parsing log files:
 - ▶ Several users may share the same IP (e.g., users behind a firewall or proxy)
 - ▶ A user may navigate with two or more open browser
 - ▶ A user may wait several minutes between sending requests, when does a session terminate?
 - ▶ A threshold needs to be defined (e.g., 30 minutes).

Using UBGs in practice

- ▶ Using UBGs, we can perform:
 - ▶ **simulation** to generate similar, but non-identical, sessions.
 - ▶ *e.g., validate the system under varying workload mixes*
 - ▶ **analysis** can help understanding user behaviour.
 - ▶ *e.g., understand how website topology affects navigation*
 - ▶ **modification** can help exploring the consequences of changes.
 - ▶ *e.g., what if we merge two pages?*
 - ▶ **clustering** can help grouping similar users into **classes**.
 - ▶ *e.g., helpful for business analytics, pre-fetching, sizing, ...*

Simulating a UBG

```
1:  $i := E$  /* initial state */
2: while simulation is not over do
3:   print  $i$ 
4:    $r :=$  a random number in  $[0,1)$ 
5:   for  $j \in \mathcal{S}$  do
6:     if  $r \leq \sum_{k \leq j} p_{ik}$  then
7:        $i = j$ 
8:       break
9:     end if
10:  end for
11: end while
```


Properties

Visit ratios:

- ▶ What is the **average number of visits** $V(i)$ to state i ? (i.e., the mean number of invocations to the corresponding page)
- ▶ In the example, $V(\text{Add}) = 0.725$ hence $V(\text{Buy}) = 0.725 \times 0.4 = 0.29$.
- ▶ Generalising the argument, we need to solve the linear system:

$$V(E) = 1, \quad V(j) = \sum_{i \in \mathcal{S}} V(i) p_{ij} \quad \forall j \in \mathcal{S} \setminus \{E\}$$

where \mathcal{S} is the set of states and $E = \text{Entry}$.

Session lengths:

- ▶ What is the **average session length** L for a user?

$$L = \sum_{i \in \mathcal{S} \setminus \{E, X\}} V(i)$$

Properties

Session length distribution:

- ▶ How often does the user reach state i after n page visits?
- ▶ $\pi^{(n)}(i)$: **probability** of being in state i at the n th invocation.
- ▶ $\pi^{(n)} = [\pi_i^{(n)} | i \in \mathcal{S}]$: **state probability vector**
- ▶ In particular, $\pi^{(0)} = [\pi_E^{(0)}, \pi_H^{(0)}, \dots, \pi_X^{(0)}] = [1, 0, \dots, 0]$
- ▶ Similarly to the visits, we need to solve the linear system:

$$\pi_j^{(n)} = \sum_{i \in \mathcal{S}} \pi_i^{(n-1)} p_{ij} \quad j \in \mathcal{S}$$


or in matrix form

$$\pi^{(n)} = \pi^{(n-1)} P \quad \Rightarrow \quad \boxed{\pi^{(n)} = \pi^{(0)} P^n}$$

Properties

Session length distribution:

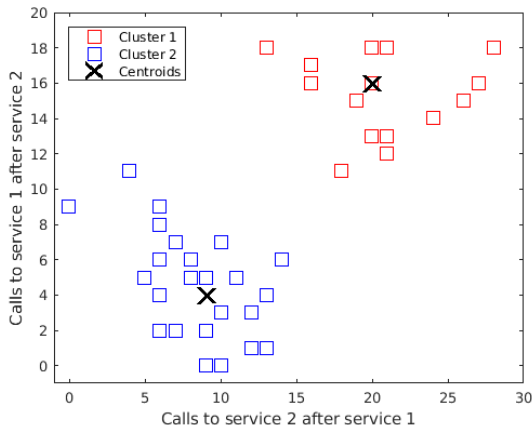
- ▶ $\pi_X^{(n)}$ probability of leaving the system in n page visits or less.
- ▶ $\pi_X^{(n)} - \pi_X^{(n-1)}$ is thus the probability of completing the session after exactly n page requests.



Fitting UBGs from log data

- ▶ UBGs can be fitted automatically from website logs
- ▶ Define for a user u the following matrix
 - ▶ C_u : entry (i, j) counts the visits to page j right after visiting i
 - ▶ After normalizing rows to sum to one, C_u becomes a UBG
- ▶ If there are n pages, C_u maps to a point in the Euclidean space with n^2 dimensions
- ▶ Each point in this space represents a possible UBG, not necessarily one observed in the logs, but one that can be easily created in load tests or other applications.

Example: centroids as typical user profiles



Each square represents a user. Each centroid models a class of users (cluster) and the associated UBG may be used for load testing, scheduling, load balancing, pre-fetching, personalization, ...

Clustering with the k -means algorithm

- ▶ k -means: a classic iterative clustering algorithm
- ▶ **Input:** points in Euclidean space, number of clusters k
- ▶ **Output:** coordinates of k centroids
- ▶ Pseudocode:
 1. Initialize centroid positions randomly
 2. Repeat until convergence of the centroid positions:
 - ▶ For every point, assign it to the cluster with nearest centroid
 - ▶ For every cluster, recalculate the position of the centroid

Clustering with the k -means algorithm

- Distances are quantified using the **Euclidean distance**

$$d(C_u, C_v) = \sqrt{\sum_{i=1}^n \sum_{j=1}^n (C_u(i, j) - C_v(i, j))^2}$$

where $C_u(i, j)$ is the element of C_u in row i and column j .


- For every cluster, centroid positions recalculated by **averaging** the coordinates of the points within the cluster

$$c(i, j) = \frac{1}{n_c} \sum_{u \in \text{cluster}} C_u(i, j)$$

where $c(i, j)$ is the centroid coordinate on dimension (i, j) and n_c counts the points currently assigned to the cluster.

Limitations of UBGs

- ▶ UBGs do not fully specify interactions, e.g., what to search?
- ▶ UBGs are agnostic of **resource usage**, e.g., CPU, memory,
- ▶ Some paths might be invalid in the real system, e.g. *Home*→*Search*→*Logout* without a *Login*.



Example: UCML - an extension of UBGs

- ▶ Different user classes, conditional actions, data flows, ...
- ▶ Other variants are integrated as tool-specific languages in load generation tools (e.g., *Jmeter*, *Load runner*, *Wessbas*, ...).

