

7. Computing on Untrusted Servers (Encrypted Databases - CryptDB)

Naranker Dulay

n.dulay@imperial.ac.uk

<https://www.doc.ic.ac.uk/~nd/peng>

Outsourcing Databases to the Cloud



- ▶ Business database needs are growing by over 50% per year.
- ▶ Costs of database management are increasing as well.
- ▶ Database outsourcing is now a viable choice for many businesses
- ▶ But should organisations trust the database service provider to protect the confidentiality of personal data (financial, medical, mail/messages etc) or company data (financial, operational, research, security, business plans etc)?

Threats

- ▶ **A cloud database provider might** breach their data processing agreement, for example:

Pass data to another party (*Theft*) or to a government agency (*Lawful Theft*)

Deny access to the database in a dispute e.g. over payment/quality of service.

Use other providers for services (e.g. storage) without permission from users.

Not have adequate security measures and gets hacked.

Keep changing terms and conditions ...

- ▶ **A privileged user (or hacker) who was malicious might be able to:**

Read (copy) data from Disk/RAM. Modify/delete data. Swap/Delay/Replay data. Manipulate access controls. Observe who accesses which data when and from where. Etc....

► **A good solution should protect against such attacks** ◀

Why not encrypt the database?

We could encrypt/decrypt all attributes at the *client* side. But what if we need a few tuples (rows) from relations (tables) with millions of tuples (rows)?

```
select * from Purchases where Bought between today-6 and today;
```

How might this be executed?

Why not encrypt the database?

We could encrypt/decrypt all attributes at the *client* side. But what if we need a few tuples (rows) from relations (tables) with millions of tuples (rows)?

```
select * from Purchases where Bought between today-6 and today;
```

We might:

- (1) **Fetch ALL** encrypted **Bought** dates and primary keys from the Server.
- (2) Decrypt **Bought** dates
- (3) Find tuples matching the **where** clause
- (4) Fetch the remaining attributes for the matched tuples from the Server using their Primary Keys
- (5) Decrypt the fetched attributes

For this particular example, we could perform a
select on 7 encrypted dates.

Encrypted Queries

Why not encrypt the data and encrypt the query?

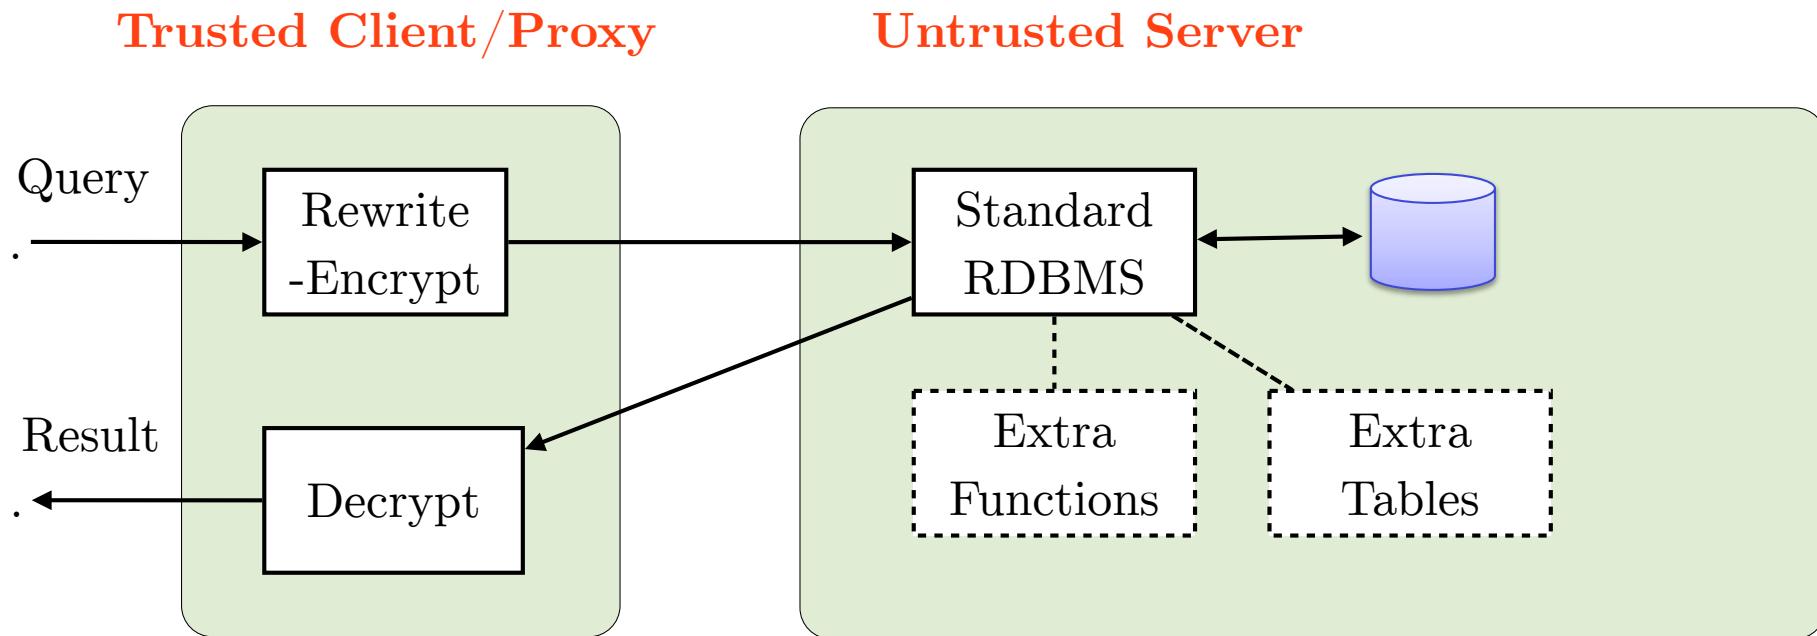
What we'd like to have, is the DBMS perform an **encrypted query** on the **encrypted data** without decrypting the data or query (c.f. FHE)

This is not easy in general since DBMSs are so feature rich.

For example a DBMS supporting encrypted processing would need to

- ▶ Handle numerous datatypes and operators — ints, strings, booleans, floats, decimals, arithmetic, comparisons, aggregate functions, sorting,
- ▶ Support transactions, access control, schema updates, etc.
- ▶ Keep memory, storage and time overheads low.
- ▶ Prevent leakage and traffic analysis, while ensuring integrity and freshness, etc.

CryptDB - Overview



In CryptDB a query like

```
select * from staff where salary >= 50000;
```

is rewritten to a query like

```
select * from table7 where col4 >= x738D3AB63;
```

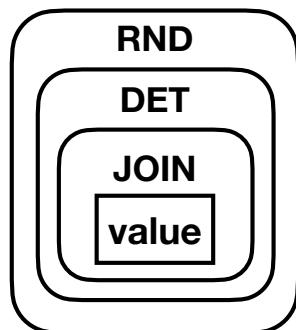
with the returned results decrypted by the client



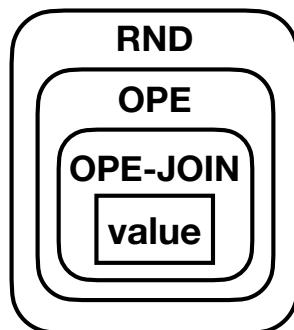
Encryption Onions

In order to provide basic SQL functionality, CryptDB employs different property-preserving encryption schemes (PPEs) and a clever approach to ‘wrapping’ and ‘unwrapping’ encryption layers.

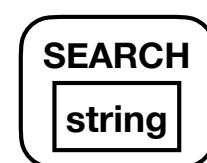
**Equals
Onion**



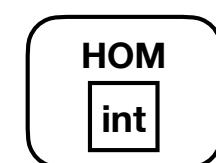
**Order
Onion**



**Search
Onion**



**Addition
Onion**



Encrypted Tables & Queries

Student		table7					
ID	NAME	c1_equals	c1_order	c1_add	c2_equals	c2_order	c2_search
32	Alice	84ab9a...	e98a..	54f3..	2445c...	abde8..	7acd..

```
update table7 set c2_equals = decrypt_RND(key, c2_equals)
```

Example

If a client requested Alice's ID, they might use:

```
select ID from Student where NAME='Alice'
```

this would be rewritten by the trusted client/proxy to

```
update table7 set c2_equals=decrypt_RND(Kt7,c2,Equals,RND, c2_equals)  
select c1_equals from table7 where c2_equals=d3ac...
```

Example continued

If the next query is

```
select COUNT(*) from Student where NAME='Bob'
```

a client can use knowledge of the previous decryption to rewrite the query to

```
select COUNT(*) from table7 where c2_equals=34bd...
```

Here 'Bob' (ciphertext 34bd...) is encrypted with the key

$K_{t7,c2,Equals,JOIN}$ and then with the key:

$K_{t7,c2,Equals,DET}$

More Encrypted Databases

- ▶ **Property-preserving encryption:**

CryptDB, Cipherbase, Monomi, ...

Google Encrypted BigQuery

Microsoft Always Encrypted v1 (included in
SQL Server 2016)

...

- ▶ **Trusted Execution Environments
(Intel SGX, FPGAs)**

EnclaveDB (Imperial), StealthDB, ObliDB,
EncDBDB,

Microsoft Always Encrypted v2 (included in
Azure SQL Database & SQL Server 2019),

...

- ▶ **Garbled circuits:**

Arx DB (Tested on ShareLaTex service
using MongoDB)



8. Computing on Untrusted Servers (Multi-user Keyword Search)

Naranker Dulay

n.dulay@imperial.ac.uk

<https://www.doc.ic.ac.uk/~nd/peng>

Multi-User Databases

CryptDB is a nice approach for many applications/scenarios but

- ▶ It was based on **secret keys shared by clients**.
- ▶ What if we need to **revoke** (change) the key, e.g. remove a user?
- ▶ We'd have to re-encrypt all the data at the database. This would be very inefficient.

Multi-User Databases

For multi-user databases we'd like to

- ▶ Allow multiple users to insert/update/search/read data
- ▶ Without sharing keys
- ▶ Easily add or revoke users without re-encrypting data.
- ▶ Use keys to authenticate users.

Shared and Searchable Data Spaces (Imperial)

- * New users can be added easily. Keys can be revoked without affecting other users
- * Keys are not shared. Keys can also be used to authenticate users.

RSA-based Proxy re-encryption scheme where

- * Data encrypted by Alice is re-encrypted by the server (proxy).
- * Re-encrypted data can be decrypted by Bob.
- * Server and Bob don't know Alice's keys.
- * Server is assumed to be **honest-but-curious** i.e. executes instructions faithfully but may be curious about what's stored or being searched for.
- * Clients are assumed to be trusted.
- * A trusted key management service generates key material for users.

RSA Proxy Key Construction

A trusted key management service constructs a master RSA Key pair (e, n) and (d, n) and then:

For each user A, the key management service constructs two new key pairs from the master RSA key pair, one pair for the user, one for the DB server:

User Key Pair 1. $Ae1 = (e1, n)$ and $Ad1 = (d1, n)$ is sent to the user

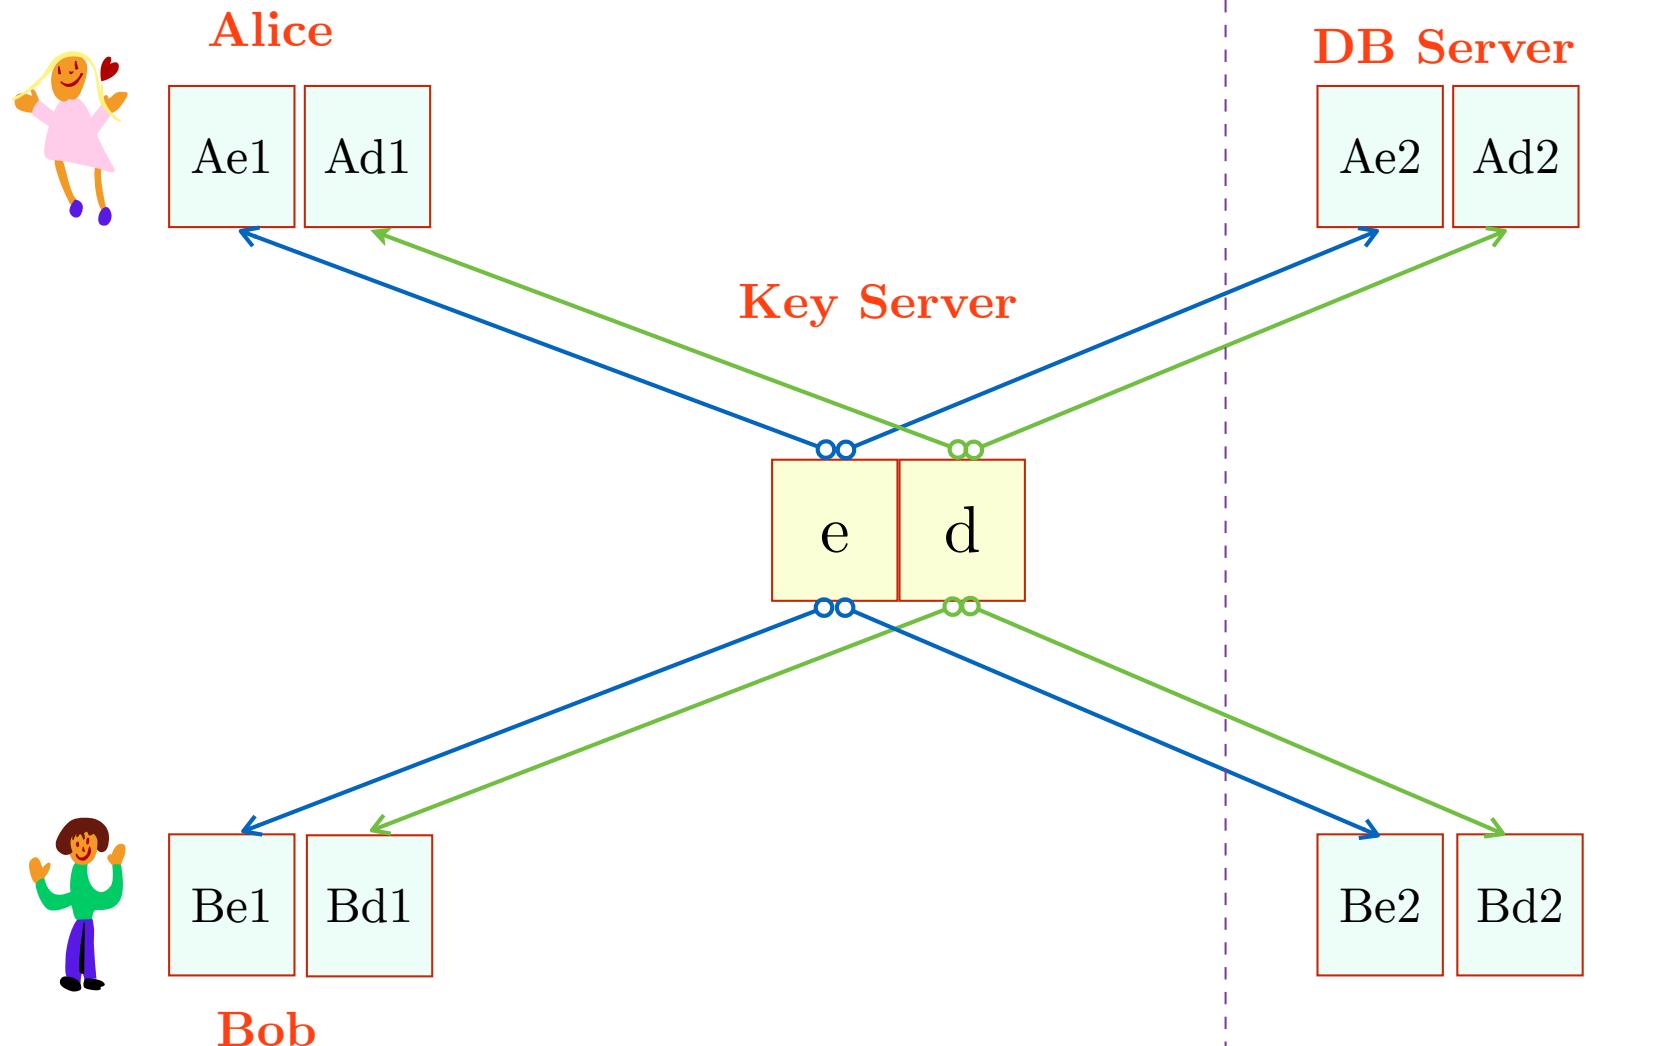
DB Key Pair 2. $Ae2 = (e2, n)$ and $Ad2 = (d2, n)$ is sent to the DB

such that

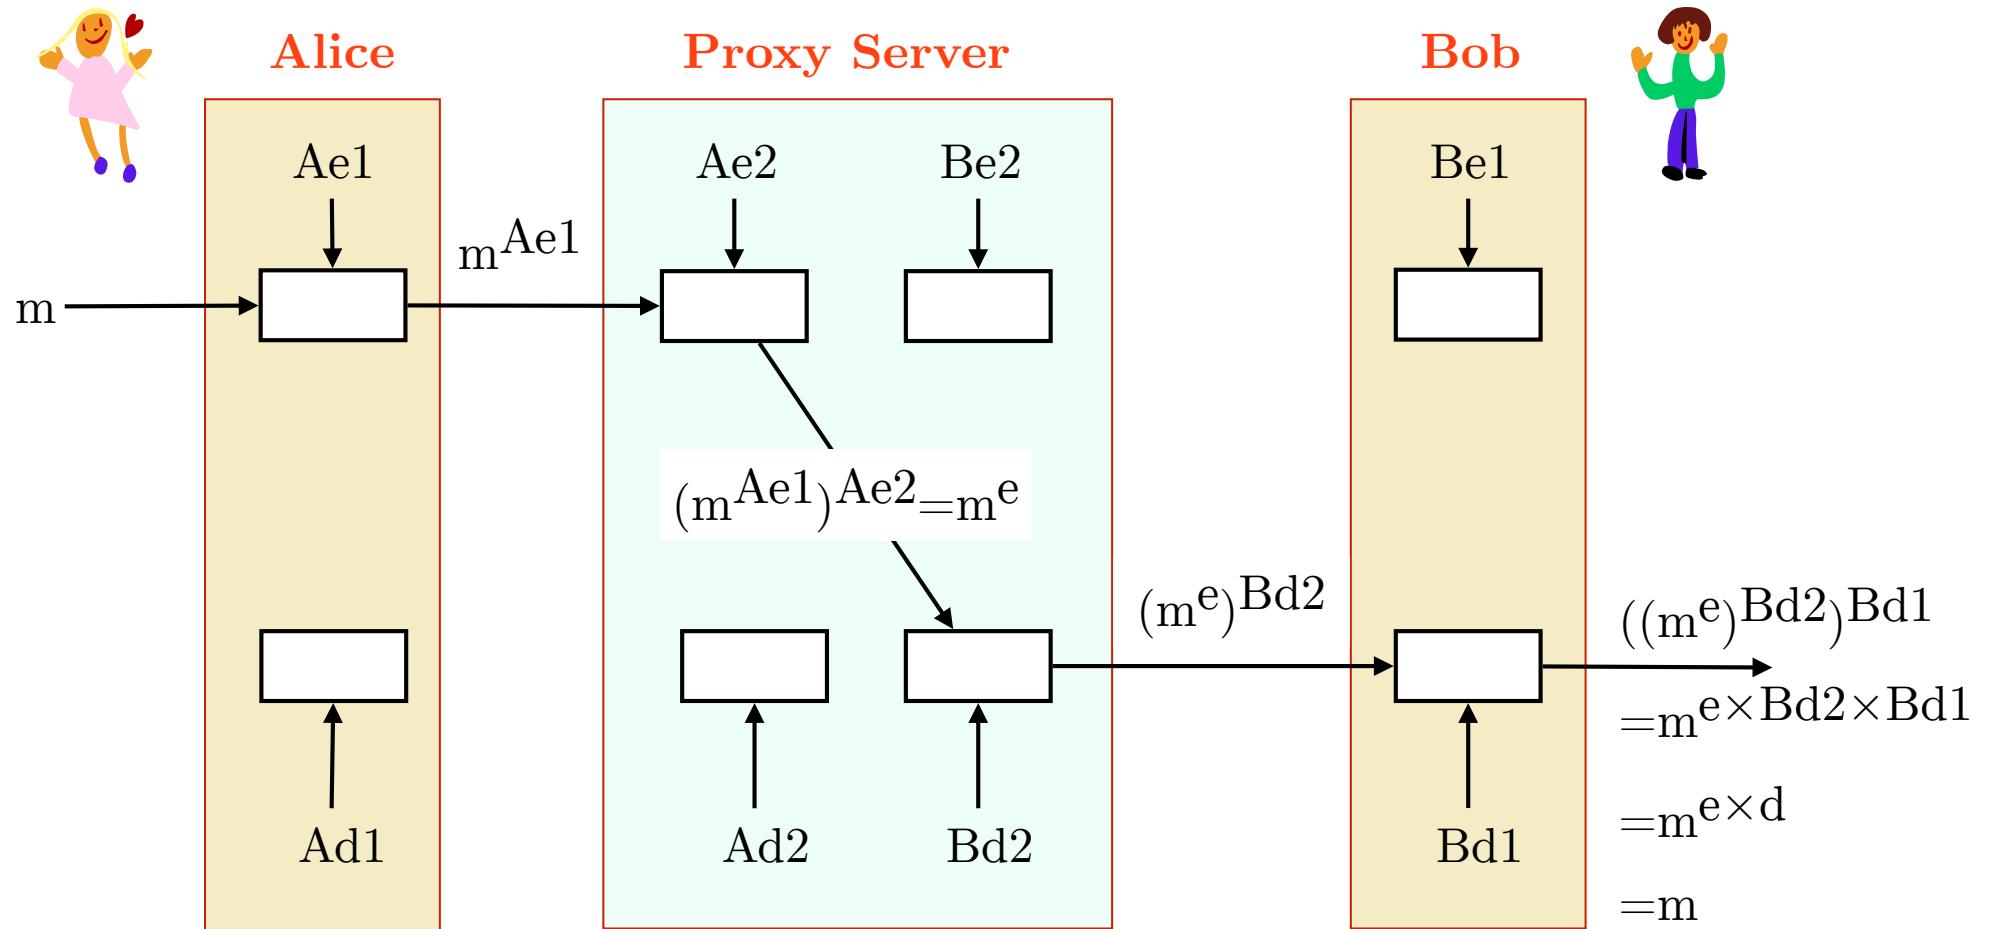
$$e1 \times e2 = e \pmod{(p-1)(q-1)}$$

$$d1 \times d2 = d \pmod{(p-1)(q-1)}$$

Key Generation



Encryption - Decryption



Example - Encrypted Keyword Search

- ▶ Alice encrypts the text to be stored with a symmetric key **k**.
- ▶ Alice then encrypts the symmetric key **k** with her public key **Ae1**
- ▶ Alice also hashes any keywords in the text and encrypts each hash value with **Ae1**
- ▶ Alice sends the (i) encrypted text, (ii) encrypted key and (iii) encrypted hash keywords to the DB server.
- ▶ DB server re-encrypts the encrypted key and encrypted hash keywords with **Ae2** (these are effectively encrypted with a master RSA encryption key **e**) and stores them, along the encrypted text (note - the encrypted text doesn't need to be re-encrypted by the DB server)
- ▶ Bob hashes his search keyword encrypts it with his public key **Be1** and sends to the DB Server as a query.
- ▶ The DB server re-encrypts Bob's encrypted hashed keyword again with **Be2**. The query is now effectively **hash(keyword)^e**.
- ▶ The DB server does a search on the query to retrieve the matching encrypted text plus its encrypted symmetric key.
- ▶ The DB server then decrypts the doubly encrypted key **k** with **Bd2** and sends it to Bob along with the encrypted text.
- ▶ Bob decrypts the key again with **Bd1** to get **k** and then uses **k** to decrypt the encrypted text.

9. Computing on Untrusted Servers (Location Sharing - Longitude)

Naranker Dulay

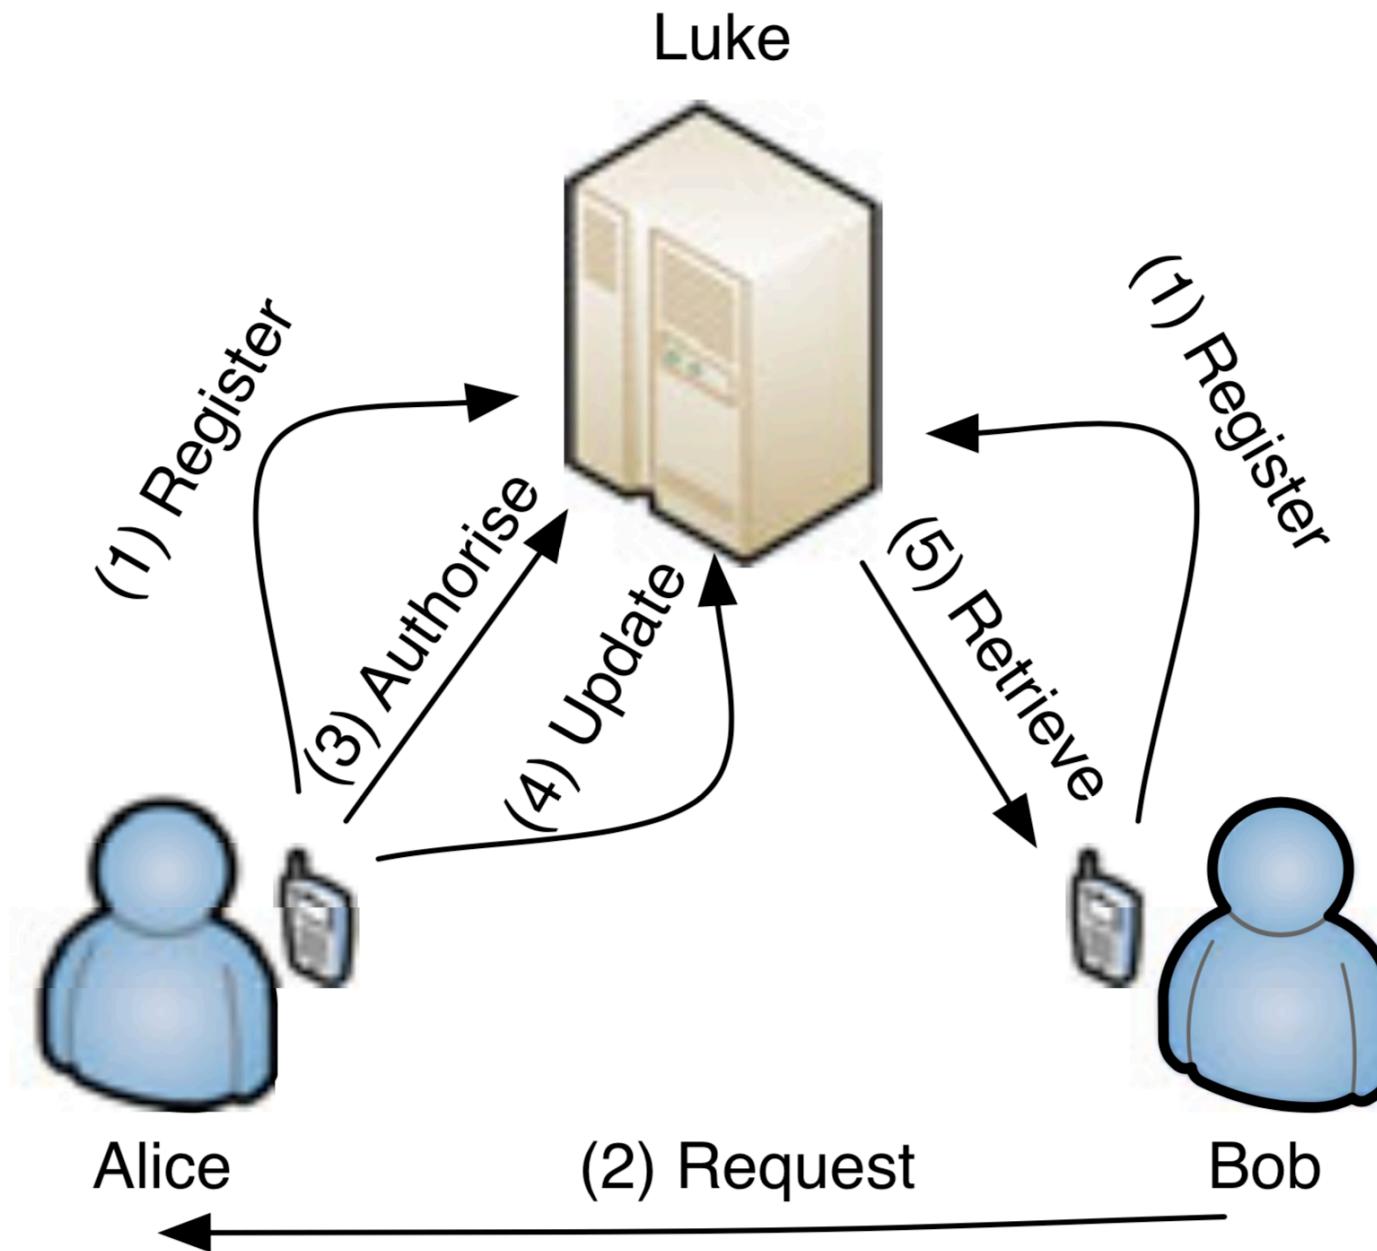
n.dulay@imperial.ac.uk

<https://www.doc.ic.ac.uk/~nd/peng>

Longitude: Privacy-preserving location sharing

- * Privacy-preserving location ***sharing*** service for **mobile applications**. c.f. Apple's Find My Friends and Google's long-abandoned Latitude service.
- * Alice shares her location with Bob via a location sharing provider Luke (who is **honest-but-curious**).
 - *What sharing policies should be supported?*
 - *What if Alice shares her location with Bob at full precision and David, her boyfriend, only at 1KM accuracy? What if David finds out?*
- * Most computational cryptographic tasks are performed by Luke (saves communication and battery). Alice's cryptographic material can be pre-computed when her phone is charging.
- * Alice only encrypts new location data once, regardless of the number of friends she shares her location with (this minimises computation and communication)

Longitude Protocol Simplified 1



Longitude Protocol Simplified 2

0. Setup

Luke: Luke generates public parameters common to all users

Alice, Bob, ...: Alice and Bob generate a public/private keypair: (\mathbf{Pa} , \mathbf{Va}), (\mathbf{Pb} , \mathbf{Vb}) respectively.

1. Registration

Alice and Bob register with Luke

Alice → Luke: RegisterMe("Alice", ...)

Bob → Luke: RegisterMe("Bob", ...)

2. Request:

Bob requests Alice to share her location with him

Bob → Alice: "Please share your location with me?", \mathbf{Pb}

Alice → Bob: "Okay"

3. Authorise:

Alice sends a **re-encryption key** to "authorise Bob at **Precision 5**" to Luke.

Alice can also set other policies like the times that sharing is permitted or not.

Alice: $\mathbf{RKab} = \text{genReEncKey}(\mathbf{Va}, \mathbf{Pb})$, $\mathbf{PRECab}=5$

Alice → Luke: Authorise("Bob", \mathbf{RKab} , \mathbf{PRECab})

4. Update:

Alice periodically sends Luke her location encrypted with *her* Public Key \mathbf{Pa} .

Luke keeps last update only.

Alice → Luke: $\mathbf{ELOCa} = \text{genEncLoc}(\mathbf{LOCa}, \mathbf{Pa})$

5. Retrieve:

Bob requests Alice's last location. Luke returns the location encrypted for Bob using Alice's Re-encryption key for Bob at the precision Alice authorised.

"Where is Alice?"

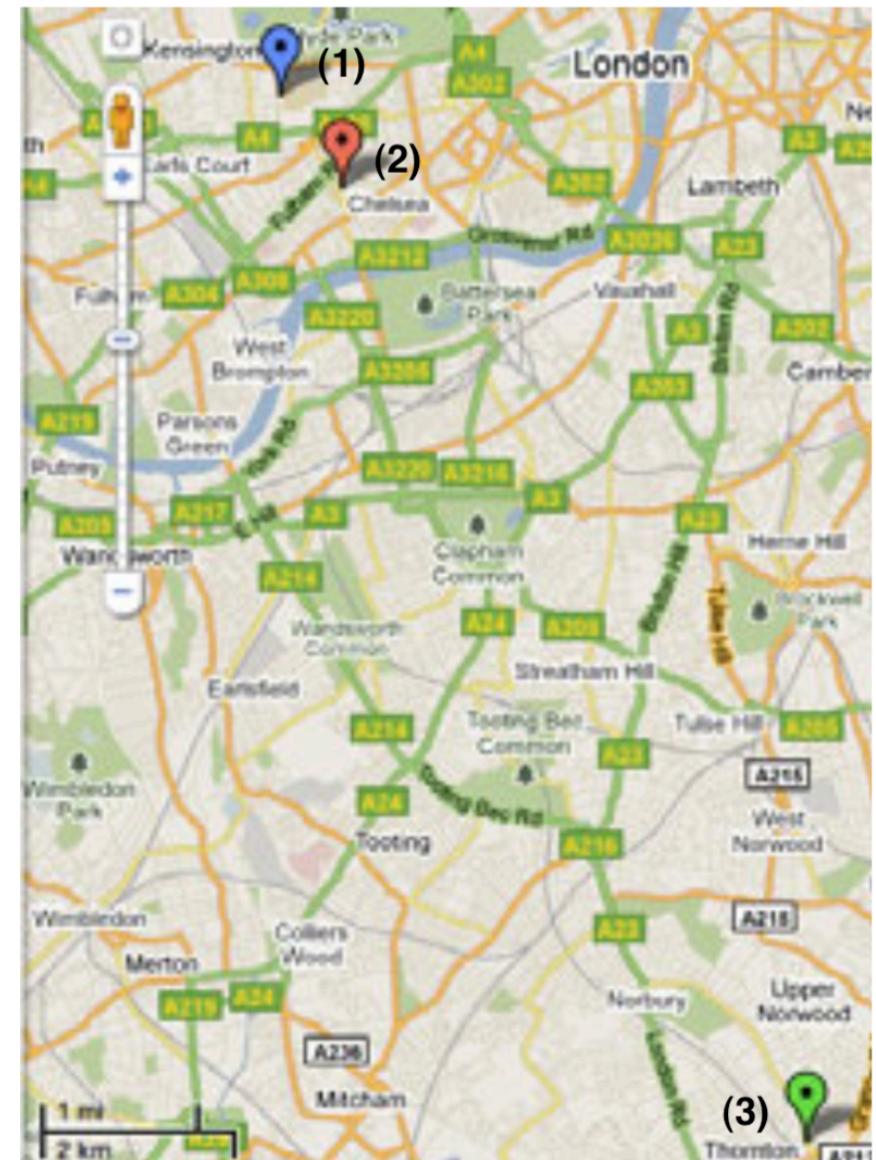
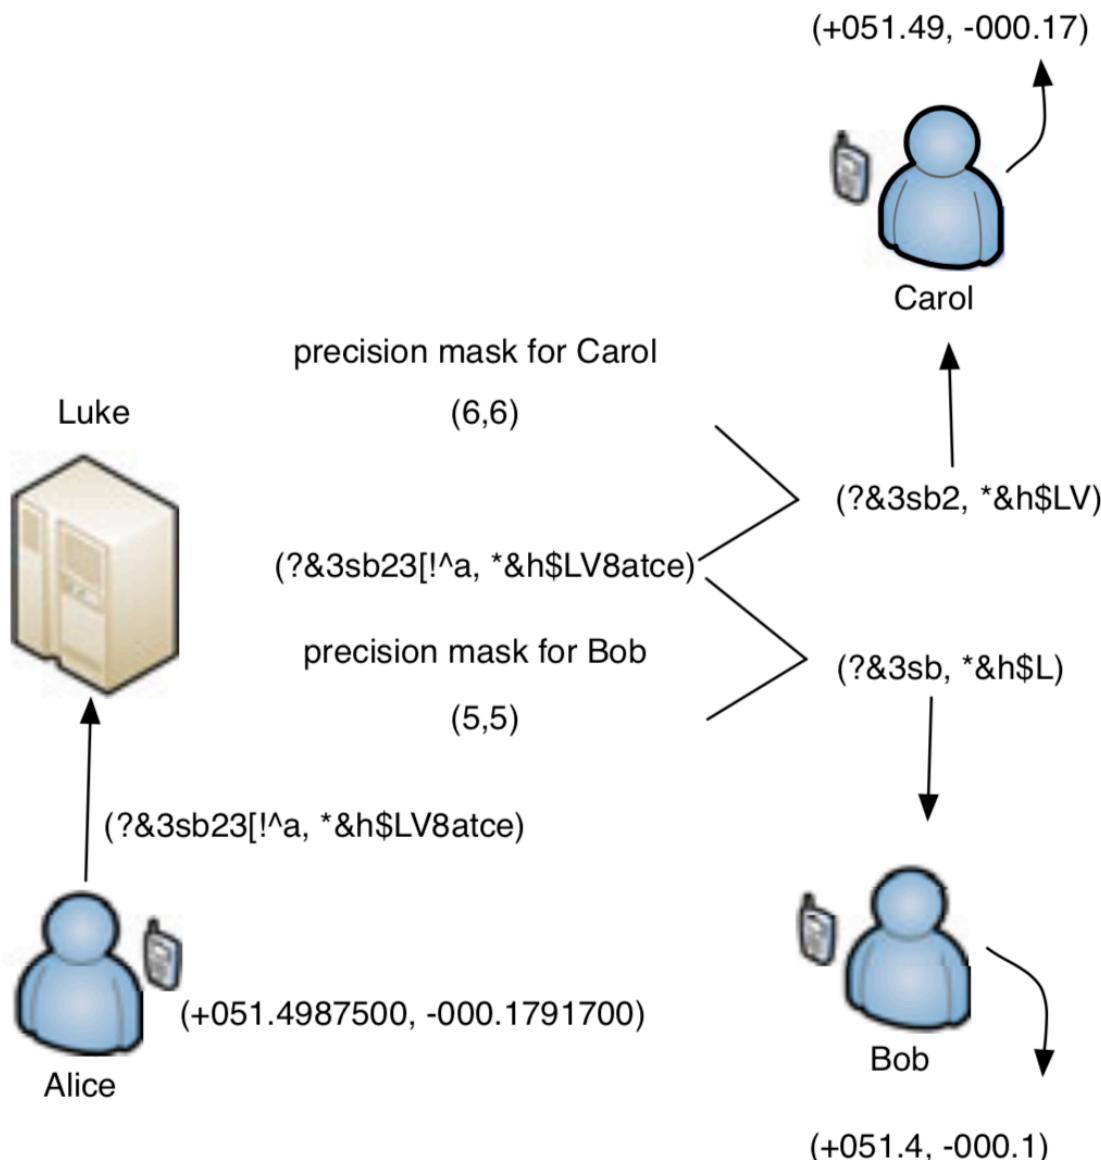
Luke → Bob: $\mathbf{ELOCab} = \text{reGenEncLoc}(\mathbf{ELOCa}, \mathbf{RKab}, \mathbf{PRECab})$

Bob: $\mathbf{LOCa} = \text{Decrypt}(\mathbf{ELOCab}, \mathbf{Vb})$

Precision of Encrypted Location

- * Latitude and Longitude coordinates are represented as a pair of decimal numbers e.g. **(51.49875, -0.17917)** for the Department of Computing.
- * 5 decimal places is precise to \sim 1m,
4 decimal places is precise to \sim 11m,
3 decimal places is precise to \sim 111m,
2 decimal places is precise to \sim 1.1km,
1 decimal places is precise to 11km.
- * Alice **encodes** each coordinate as a 11 character string **SIIIFFFFFFF** (Sign, integer, fraction) e.g. 51.49875 would be encoded as "**+0514987500**"
- * This is converted into bits and xor'ed with a random bit stream.
- * The encrypted coordinate is truncated by Luke when sending location to Bob respecting Alice's policy - the precision level set by Alice in step 4. Luke does not need to decrypt the data to do this, he does it "**blindly**".

Applying a precision mask to encrypted location



Friend Revocation

- ***How can Alice stop Bob from accessing her location (revocation)?***

Alice could ask Luke not to send any further updates to Bob, i.e. Luke removes the AliceBob Re-Encryption Key (RKab). Note: RKab is not the same as RKba .

- ***But what if Luke colludes with Bob?***

To prevent this Alice can generate a new public-private keypair using parts of the old keypair and new re-encryption keys using parts of the old re-encryption keys for her remaining location-sharing friends (see tutorial).

- ***Also will Alice's remaining friends need to be notified and will they need to generate new re-encryption keys for Alice since the re-encryption key for Alice is generated from the friend's private key and Alice's public key?***

Not in Longitude. In Longitude the friends re-encryption key for Alice is generated using an unchanged part of Alice's public key.

Proxy Re-encryption in Longitude

- * Longitude's proxy re-encryption scheme is based on symmetric **bilinear pairings** that pair two elements from one *group* to an element of a second *group*.
- * \mathbf{G}_1 and \mathbf{G}_2 , are two cyclic groups of prime order q . g is a generator for \mathbf{G}_1 (See Smart p4-p5 for a simple description of cyclic groups and generators).

The bilinear pairing $e: \mathbf{G}_1 \times \mathbf{G}_1 \rightarrow \mathbf{G}_2$ has the following properties:

Bilinearity:

forall $u, v \in \mathbf{G}_1, a, b \in \mathbb{Z}q$ we have $e(u^a, v^b) = e(u, v)^{ab}$

Non-degeneracy:

$$e(g, g) \neq 1$$

Computable:

There exists an efficient algorithm to compute $e(u, v)$ forall $u, v \in \mathbf{G}_1$

1

Public parameters (G_1, G_2, e, g)



3

Alice's re-encryption key for Bob:

$$rk_{a \rightarrow b} = (h_{b1}^n, g^n \cdot h_{a2}^{-x_a})$$

$n \in \mathbb{Z}_q$ random number

2



Alice

Alice generates 3 random numbers

$x_a, y_a, z_a \in \mathbb{Z}_q$ plus her public key:

$$h_{a1} = g^{y_a}$$

$$h_{a2} = g^{z_a}$$

$$Z_a = e(g^{x_a}, g^{z_a}) = e(g, g)^{x_a z_a}$$

$$pk_a = (h_{a1}, h_{a2}, Z_a)$$

and private (secret) key:

$$sk_a = (x_a, y_a)$$

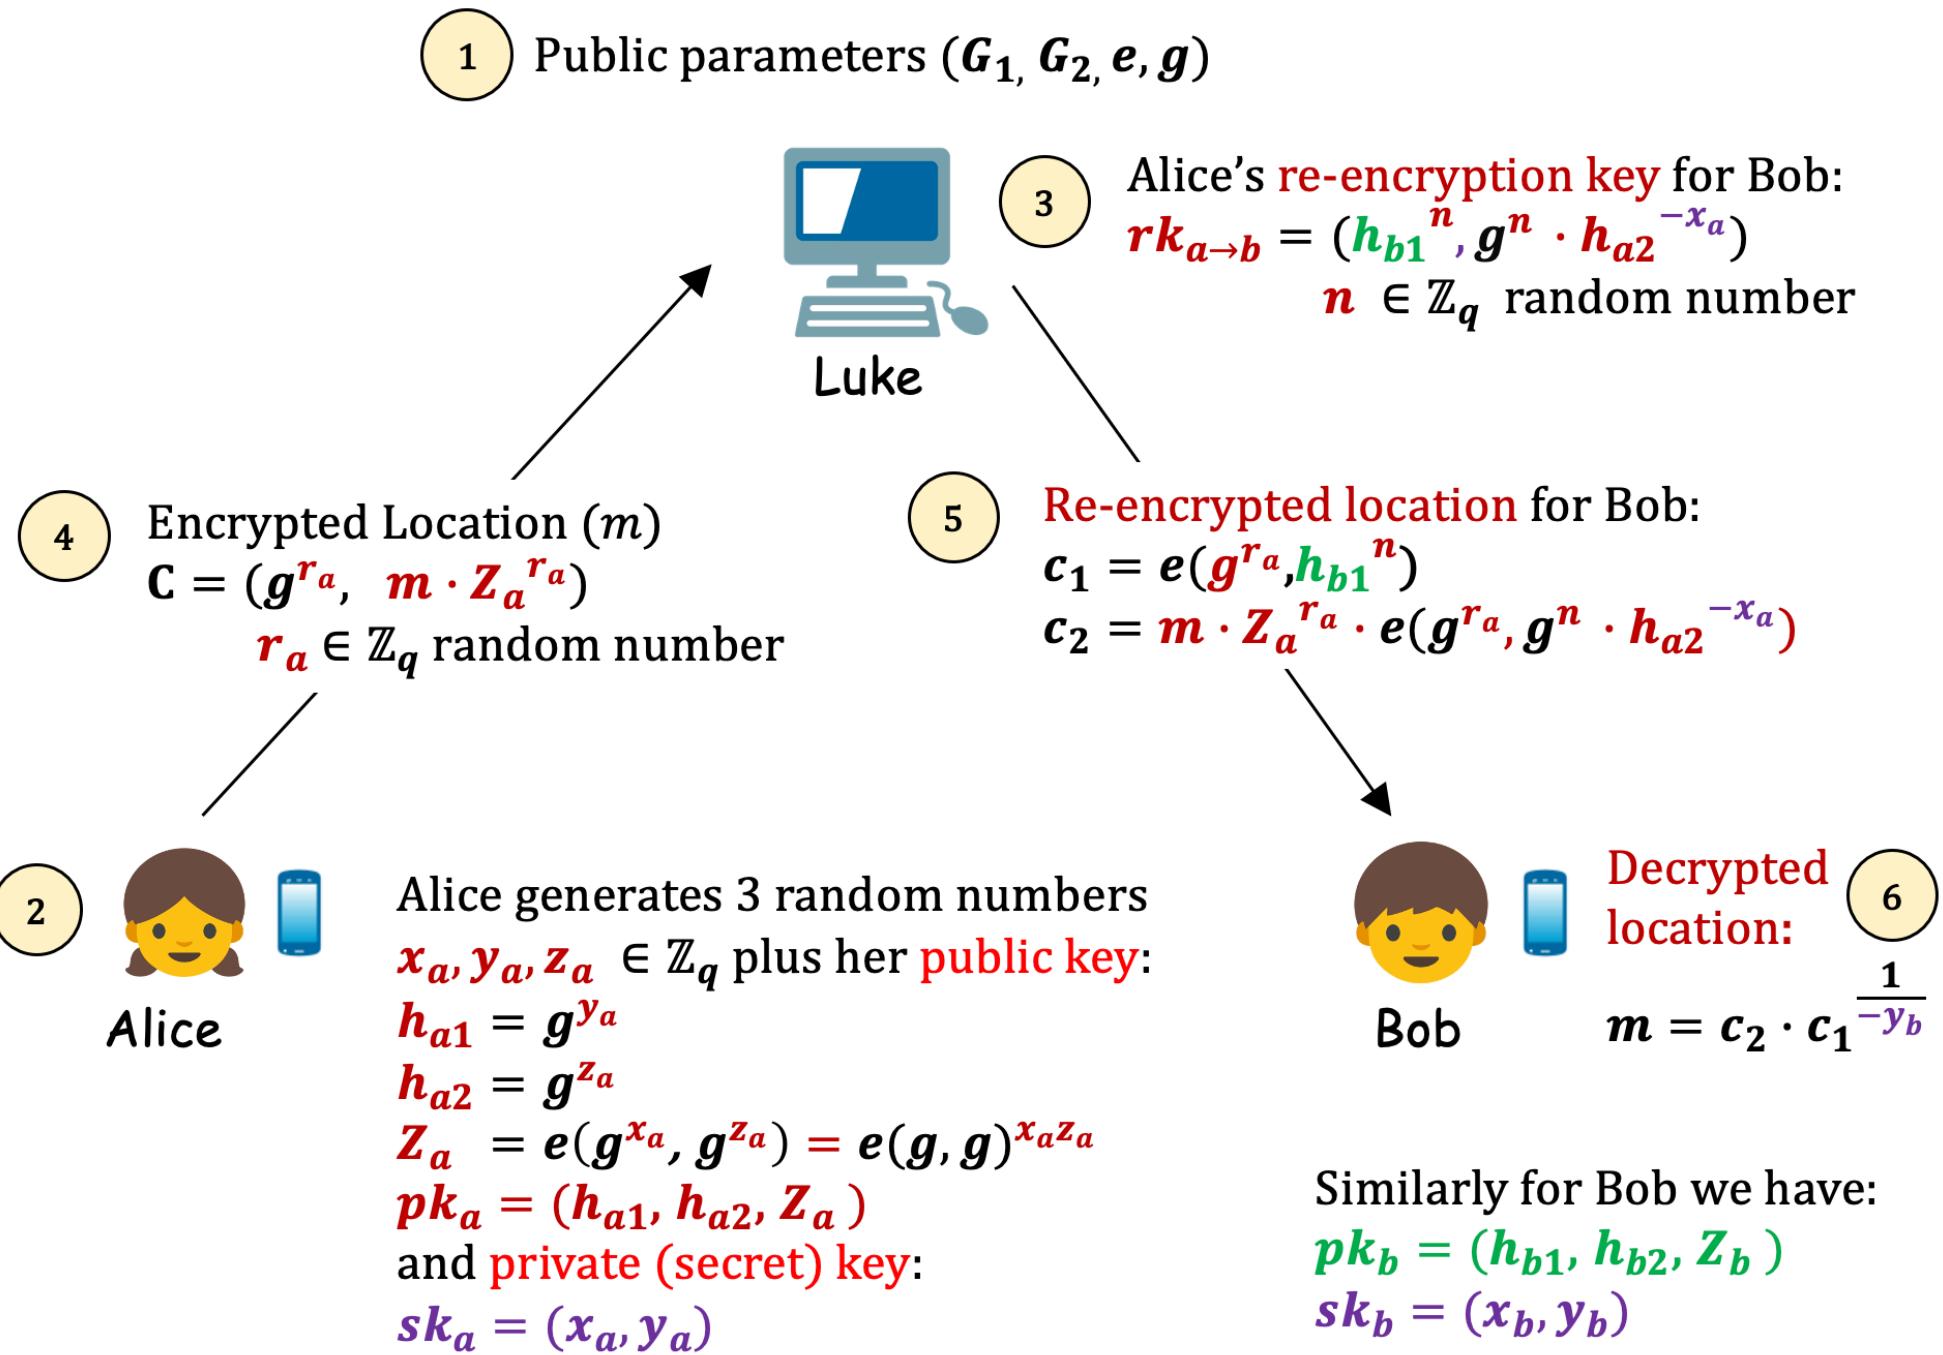


Bob

Similarly for Bob we have:

$$pk_b = (h_{b1}, h_{b2}, Z_b)$$

$$sk_b = (x_b, y_b)$$



Private Clouds

