

In-Memory Databases

OLAP versus OLTP


- OLAP (Business Intelligence):
 - Massive amounts of data
 - Complex queries
 - Large number of tables
 - Long running but still somewhat interactive
 - Mostly read only
- OLTP
 - Really only transactions, i.e., updates
 - Few tables touched
 - Typically generated queries

Reality Check - Size

- Transactional database size grows at the rate transactions increase
- 1 Tbyte is a really big TP database
- 1 Tbyte of main memory buyable for around \$50K
 - (say) 64 Gbytes per server in 16 servers
- I.e. Moore's law has eclipsed TP data base size
- If your data doesn't fit in main memory now, then wait a couple of years and it will.....

Reality Check - Performance

- TPC-C CPU cycles
- On the Shore DBMS prototype
- OldSQL should be similar



To Go Faster...

- Focus on overhead
 - Better B-trees affects only 4% of the path length
- Get rid of ALL major sources of overhead
 - Main memory deployment – gets rid of buffer pool
 - Leaving other 75% of overhead intact
 - i.e., win is 25%

Solution Choices

- OldSQL
 - Legacy RDBMS vendors
- NoSQL
 - Give up SQL and ACID for performance
- NewSQL
 - Preserve SQL and ACID
 - Get performance from a new architecture

OldSQL

Traditional SQL vendors (the “elephants”)

- Code lines dating from the 1980' s
- “bloatware”
- Mediocre performance on New TP

NoSQL

- Give up SQL
- Give up ACID

Give Up SQL?

- Compiler translates SQL at compile time into a sequence of low level operations
- Similar to what the NoSQL products make you program in your application
- 30 years of RDBMS experience
 - Hard to beat the compiler
 - High level languages are good (data independence, less code)
 - Stored procedures are good!
 - One round trip from app to DBMS rather than one one round trip per record
 - Move the code to the data, not the other way around

Give Up ACID

- If you need data consistency, giving up ACID is a decision to tear your hair out by doing database “heavy lifting” in user code
- Can you guarantee you won’t need ACID tomorrow?

ACID = goodness, in spite of what noSQL guys say

Who Needs ACID?

- Funds transfer
 - Or anybody moving something from X to Y
- Anybody with integrity constraints
 - Back out if fails
 - Anybody for whom “usually ships in 24 hours” is not an acceptable outcome
- Anybody with a multi-record state
 - E.g. move and shoot

NoSQL Summary

- Appropriate for non-transactional systems
- Appropriate for single record transactions that are commutative
- Not a good fit for New TP
- Use the right tool for the job
- But: Two recently-proposed NoSQL language standards – CQL and UnQL – are amazingly similar to (you guessed it!) SQL

NewSQL

- SQL
- ACID
- Performance and scalability through modern innovative software architecture

NewSQL

- Needs something other than traditional record level locking (1st big source of overhead)
 - timestamp order
 - MVCC
- Needs a solution to buffer pool overhead (2nd big source of overhead)
 - Main memory (at least for data that is not cold)
 - Some other way to reduce buffer pool cost
- Needs a solution to latching for shared data structures (3rd big source of overhead)
 - Some innovative use of B-trees
 - Single-threading

NewSQL


- Needs a solution to write-ahead logging (4th big source of overhead)
 - Obvious answer is built-in replication and failover
 - New TP views this as a requirement anyway

A NewSQL Example – VoltDB


- Main-memory storage
- Single threaded, run transaction to completion
 - No locking
 - No latching
- Built-in high availability and durability
 - No log (in the traditional sense)

Where all the time goes... revisited

Before



VoltDB



Current VoltDB Status

- Runs a subset of SQL
- On VoltDB clusters (in memory on commodity hardware)
- With LAN and WAN replication
- 70x a popular OldSQL DBMS on TPC-C
- 5-7x Cassandra on VoltDB key-value layer
- Scales to 384 cores


Summary

Old TP



New TP



| | | |
|---------------------|---|--|
| OldSQL for New OLTP |  | <ul style="list-style-type: none">▪ Too slow▪ Does not scale |
| NoSQL for New OLTP |  | <ul style="list-style-type: none">▪ Lacks consistency guarantees▪ Low-level interface |
| NewSQL for New OLTP |  | <ul style="list-style-type: none">▪ Fast, scalable and consistent▪ Supports SQL |

Technical Overview


“OLTP Through the Looking Glass”

<http://cs-www.cs.yale.edu/homes/dna/papers/oltpperf-sigmod08.pdf>

VoltDB avoids the overhead of traditional databases

- K-safety for fault tolerance
- no logging
- In memory operation for maximum throughput
- no buffer management
- Partitions operate autonomously and single-threaded
- no latching or locking

Built to horizontally scale



Technical Overview – Partitions (1/3)

One partition per physical CPU core

- Each physical server has multiple VoltDB partitions

Data - Two types of tables

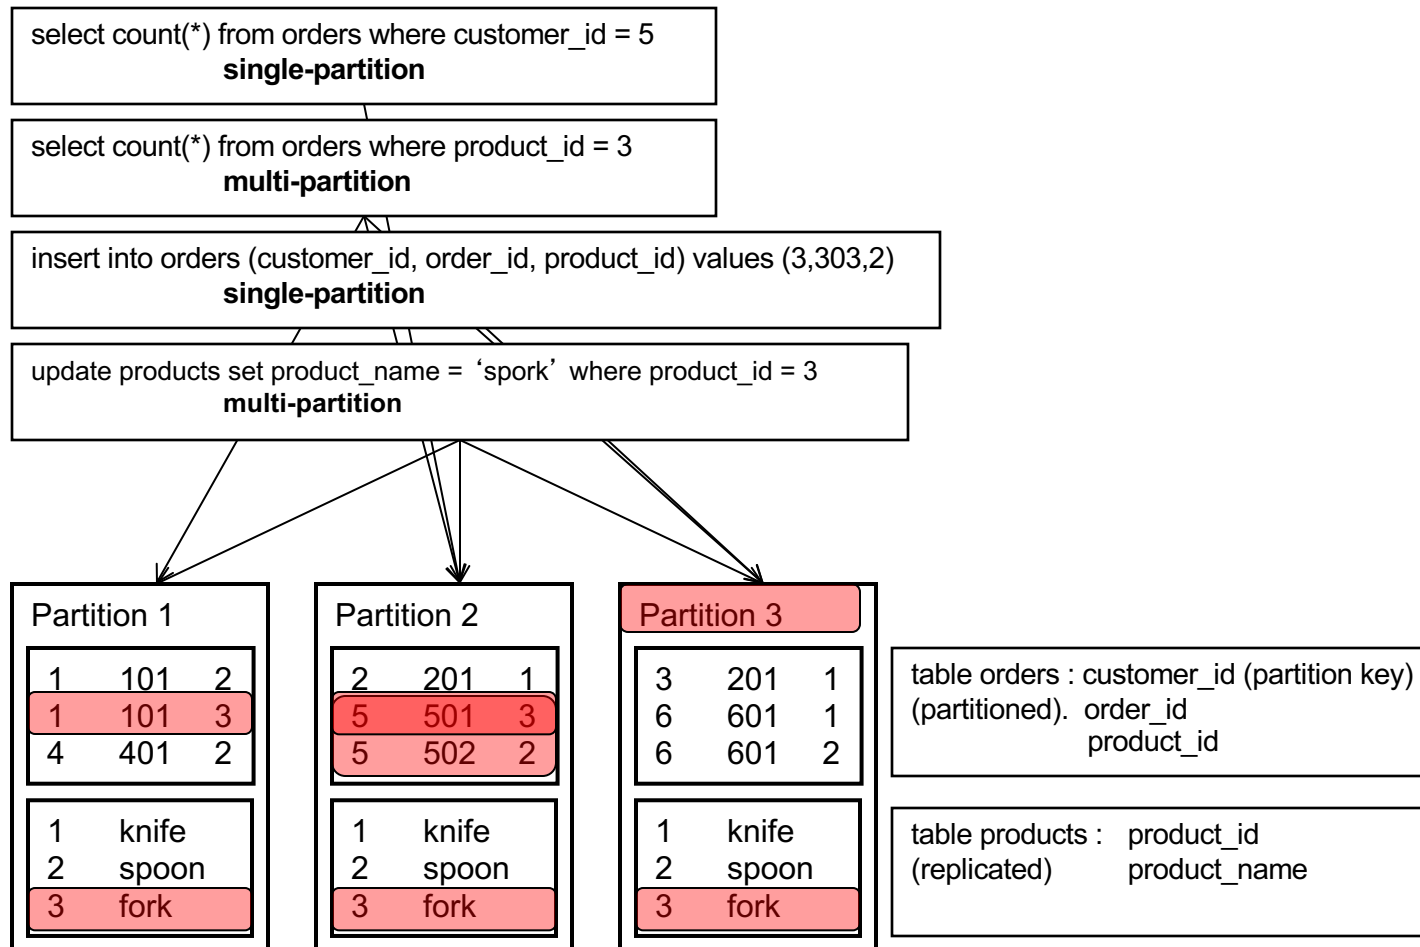
- Partitioned
 - Single column serves as partitioning key
 - Rows are spread across all VoltDB partitions by partition column
 - Transactional data (high frequency of modification)
- Replicated
 - All rows exist within all VoltDB partitions
 - Relatively static data (low frequency of modification)

Code - Two types of work – both ACID

- Single-Partition
 - All insert/update/delete operations within single partition
 - Majority of transactional workload
- Multi-Partition
 - CRUD against partitioned tables across multiple partitions
 - Insert/update/delete on replicated tables

Technical Overview – Partitions (2/3)

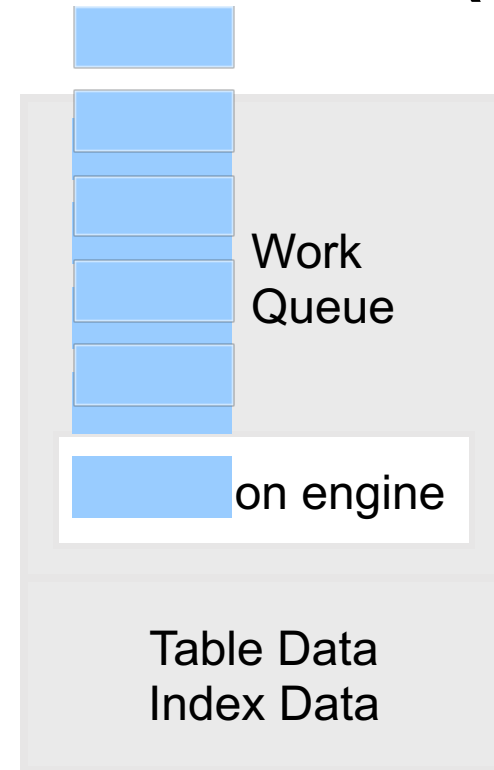
Single-partition vs. Multi-partition



Technical Overview – Partitions (3/3)

Inside a VoltDB partition...

- Each partition contains data and an execution engine.
- The execution engine contains a queue for transaction requests.
- Requests are executed sequentially (single threaded).



- Complete copy of all replicated tables
- Portion of rows (about 1/partitions) of all partitioned tables

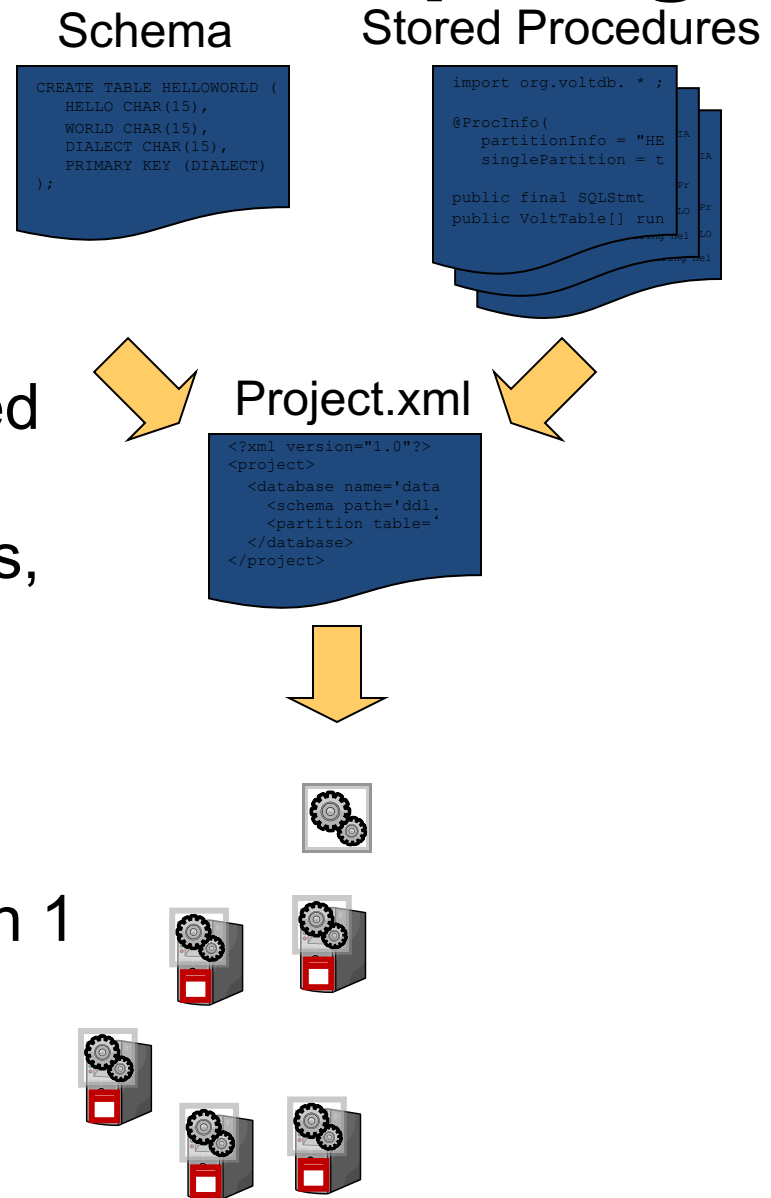
Technical Overview – Compiling

The database is constructed from

- The schema (DDL)
- The work load (Java stored procedures)
- The Project (users, groups, partitioning)

VoltCompiler creates application catalog

- Copy to servers along with 1 .jar and 1 .so
- Start servers



Technical Overview - Transactions

- All access to VoltDB is via Java stored procedures (Java + SQL)
- A single invocation of a stored procedure is a transaction (committed on success)
- Limits round trips between DBMS and application
- High performance client applications communicate asynchronously with VoltDB



Technical Overview – Clusters/Durability

- Scalability
 - Increase RAM in servers to add capacity
 - Add servers to increase performance / capacity
 - Consistently measuring 90% of single-node performance increase per additional node
- High availability
 - K-safety for redundancy
- Snapshots
 - Scheduled, continuous, on demand
- Spooling to data warehouse
- Disaster Recovery/WAN replication (Future)
 - Asynchronous replication

VoltDB and OLTP

Asynchronous Communications

Client applications communicate asynchronously with VoltDB

- Stored procedure invocations are placed “on the wire”
- Responses are pulled from the server
- Allows a single client application to generate > 100K TPS
- Client library will simulate synchronous if needed

Traditional

```
salary := get_salary(employee_id);
```

VoltDB

```
callProcedure(asyncCallback, “get_salary”, employee_id);
```

Transaction Control

VoltDB does not support client-side transaction control

- Client applications cannot:
 - `insert into t_colors (color_name) values ('purple');`
 - `rollback;`
- Stored procedures commit if successful, rollback if failed
- Client code in stored procedure can call for rollback

Lack of concurrency

- Single-threaded execution within partitions (single-partition) or across partitions (multi-partition)
- No need to worry about locking/dead-locks
 - great for “inventory” type applications
 - checking inventory levels
 - creating line items for customers
- Because of this, transactions execute in microseconds
- However, single-threaded comes at a price
 - Other transactions wait for running transaction to complete
 - Don't do anything crazy in a SP (request web page, send email)
 - Useful for OLTP, not OLAP

Throughput vs. Latency

- VoltDB is built for throughput over latency
- Latency measured in mid single-digits in a properly sized cluster
- Do not estimate latency as $(1 / \text{TPS})$

SQL Support

- SELECT, INSERT (using values), UPDATE, and DELETE
- Aggregate SQL supports AVG, COUNT, MAX, MIN, SUM
- Materialized views using COUNT and SUM
- Hash and Tree Indexes
- SQL functions and functionality will be added over time, for now done in Java
- Execution plan for all SQL is created at compile time and available for analysis

Schema Changes

Traditional OLTP

- add table...
- alter table...

VoltDB

- modify schema and stored procedures
- build catalog
- deploy catalog