

Graph Databases

Graph Databases

- Data Model:
 - Nodes and Relationships
- Examples:
 - Neo4j, OrientDB, InfiniteGraph, AllegroGraph

Graph Databases: Pros and Cons

- Pros:
 - Powerful data model, as general as RDBMS
 - Connected data locally indexed
 - Easy to query
 - Scales up reasonably well
- Cons
 - Sharding (lots of people working on this)

Introduction to Graphs

What are graphs good for?

- Recommendations
- Business intelligence
- Social computing
- Geospatial
- Systems management
- Web of things
- Genealogy
- Time series data
- Product catalogue
- Web analytics
- Scientific computing (especially bioinformatics)
- And much more!

What is a Graph?

An abstract representation of a set of objects where some pairs are connected by links.



Object (Vertex, Node)



Link (Edge, Arc,
Relationship)

Different Kinds of Graphs

- Undirected Graph



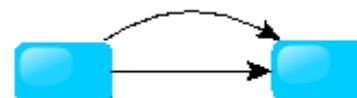
- Directed Graph



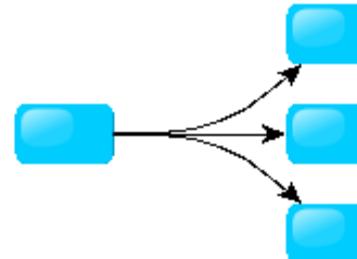
- Pseudo Graph



- Multi Graph

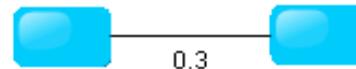


- Hyper Graph

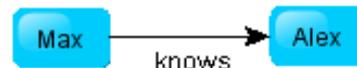


More Kinds of Graphs

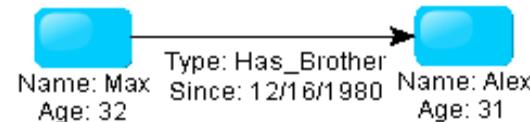
- Weighted Graph



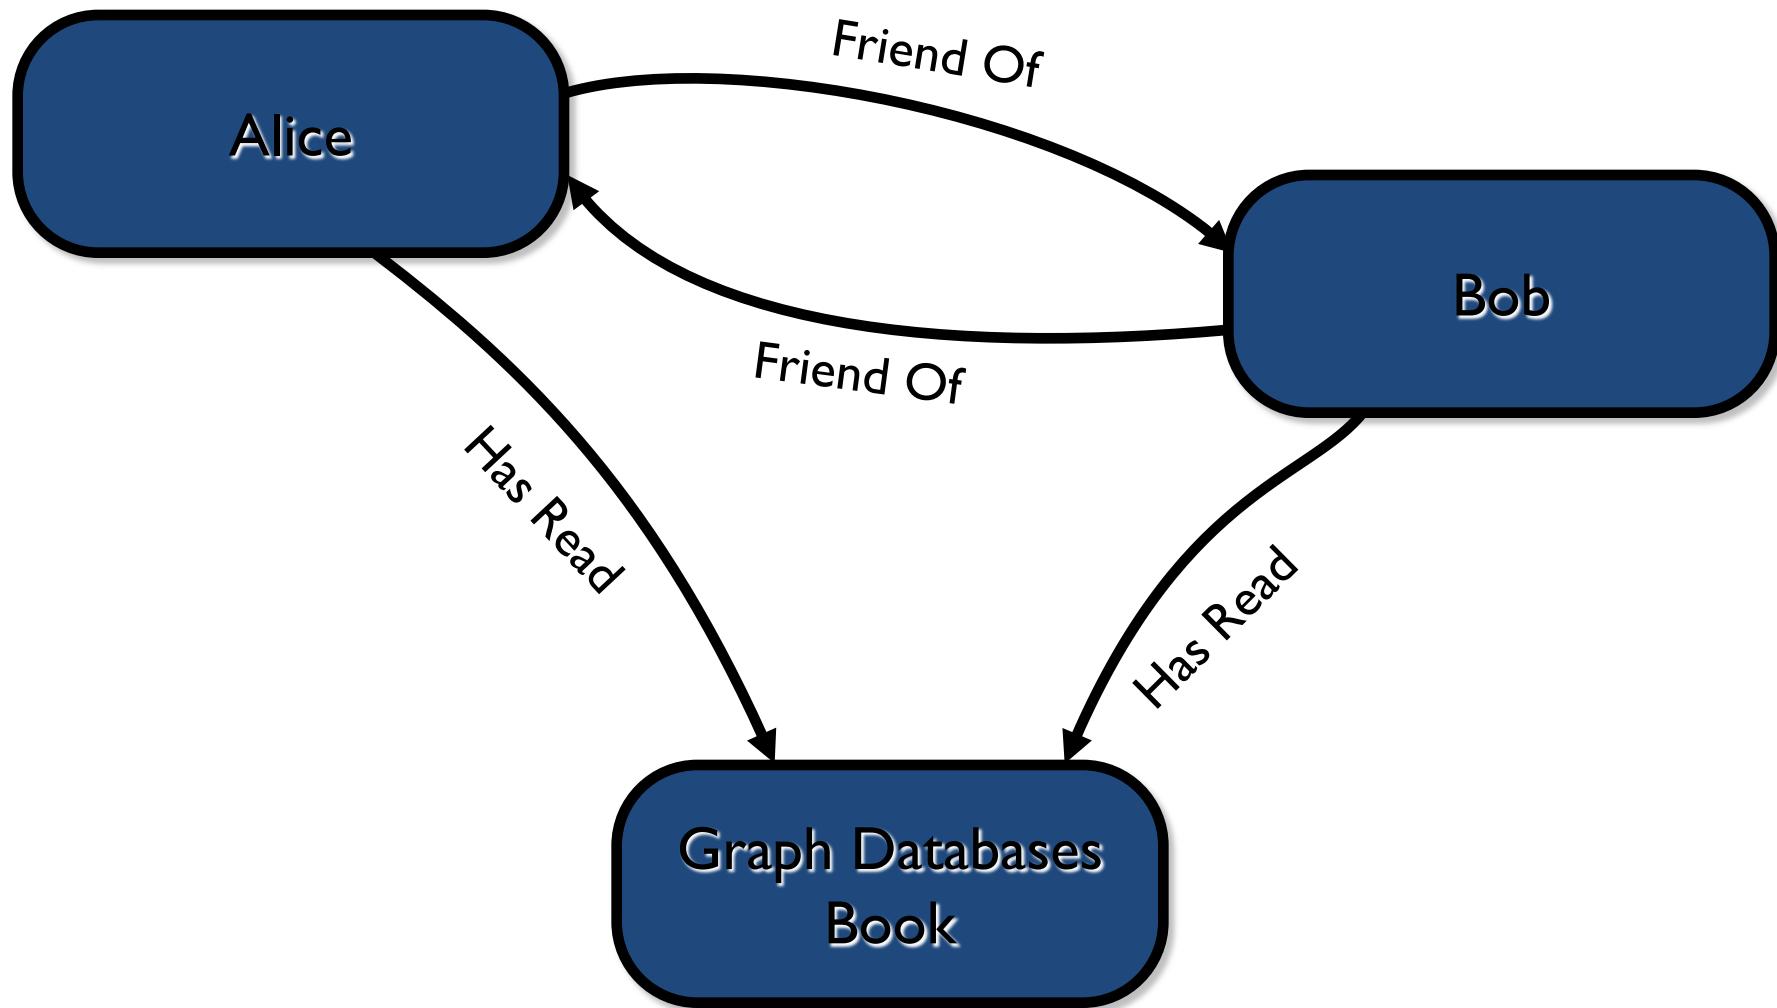
- Labeled Graph

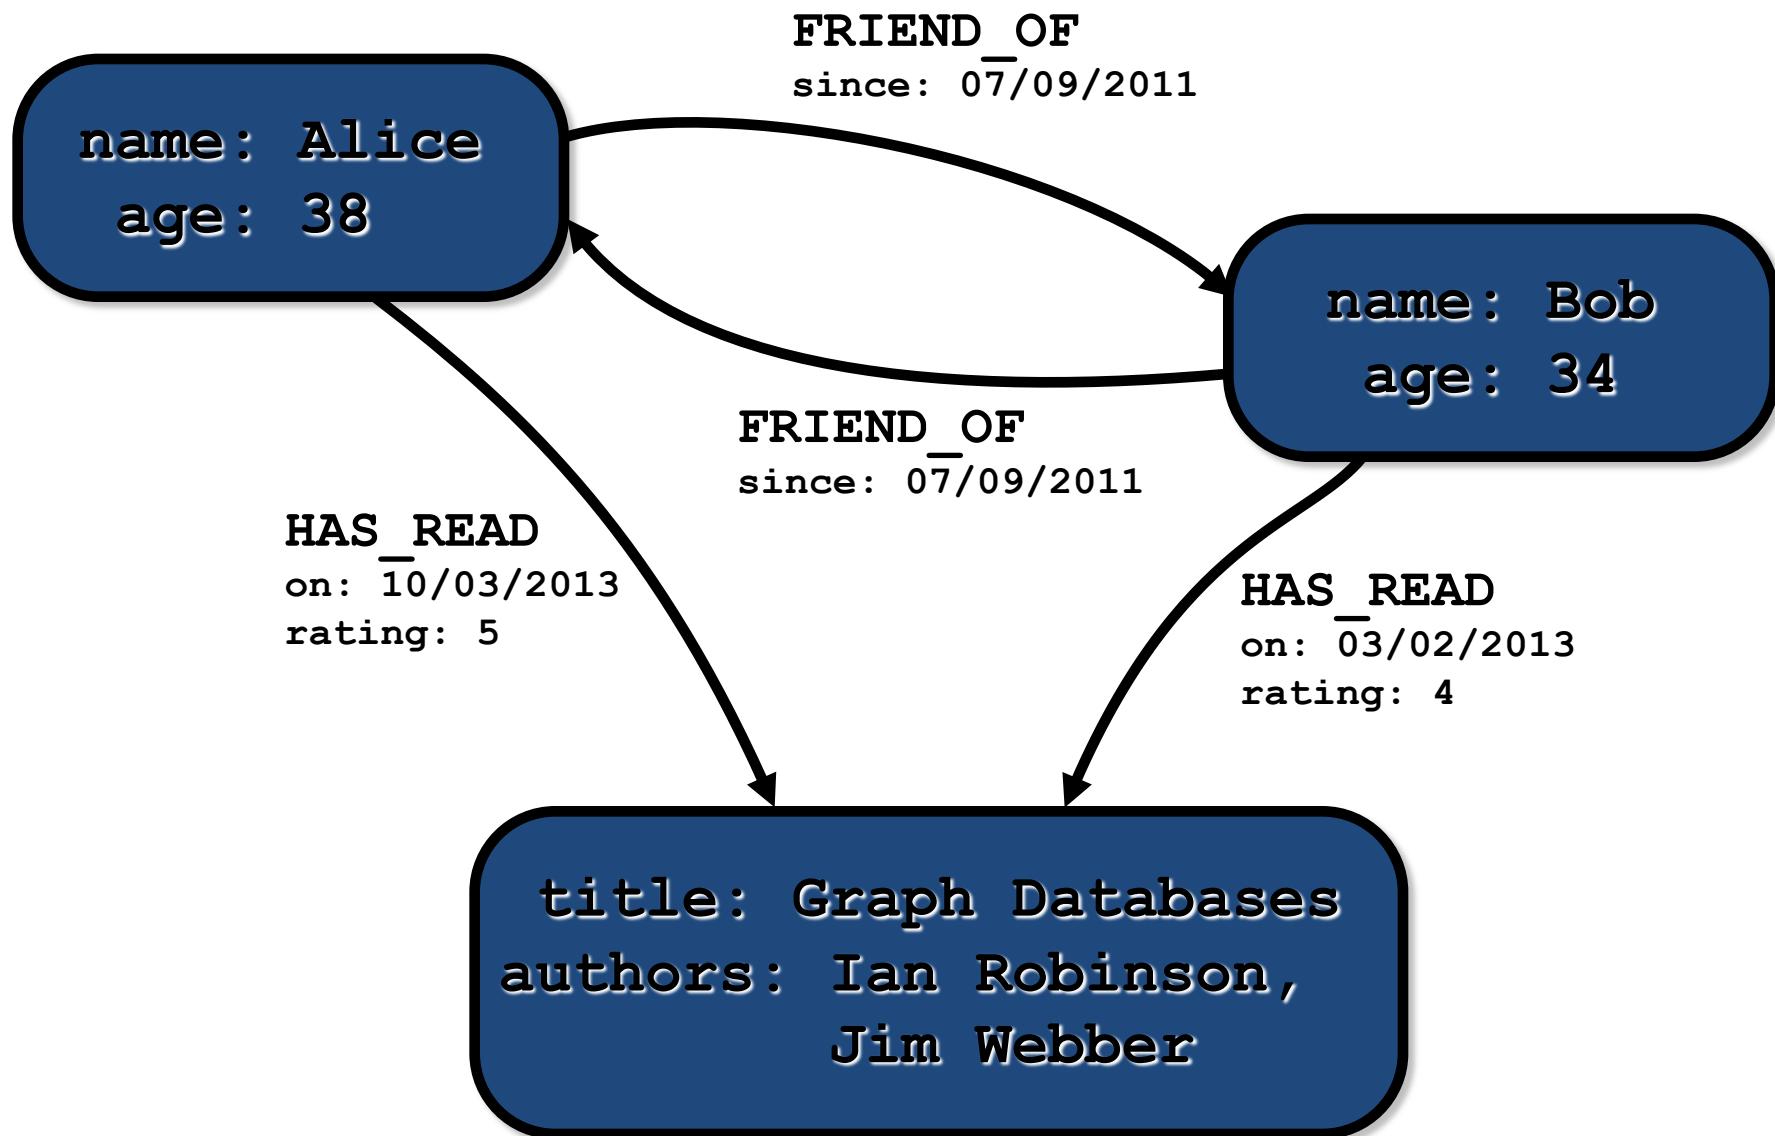


- Property Graph



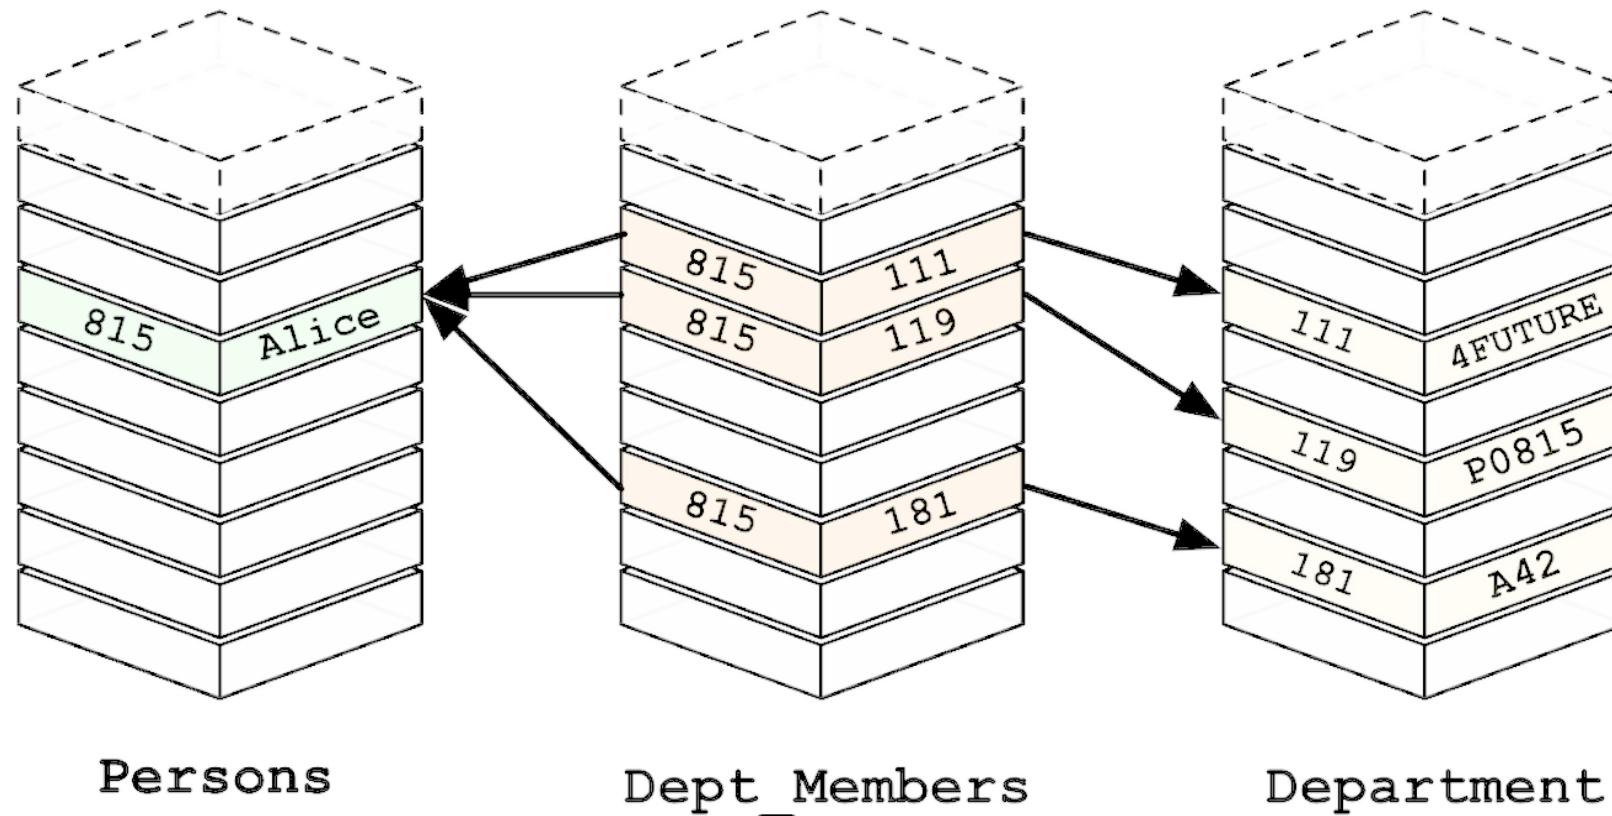
Data in Graph Databases



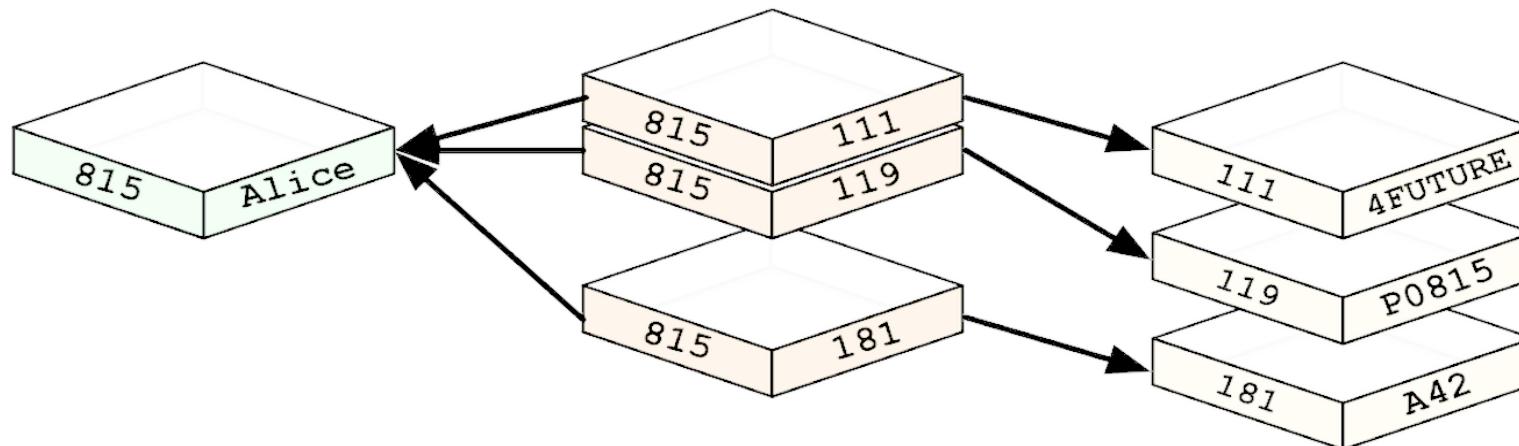


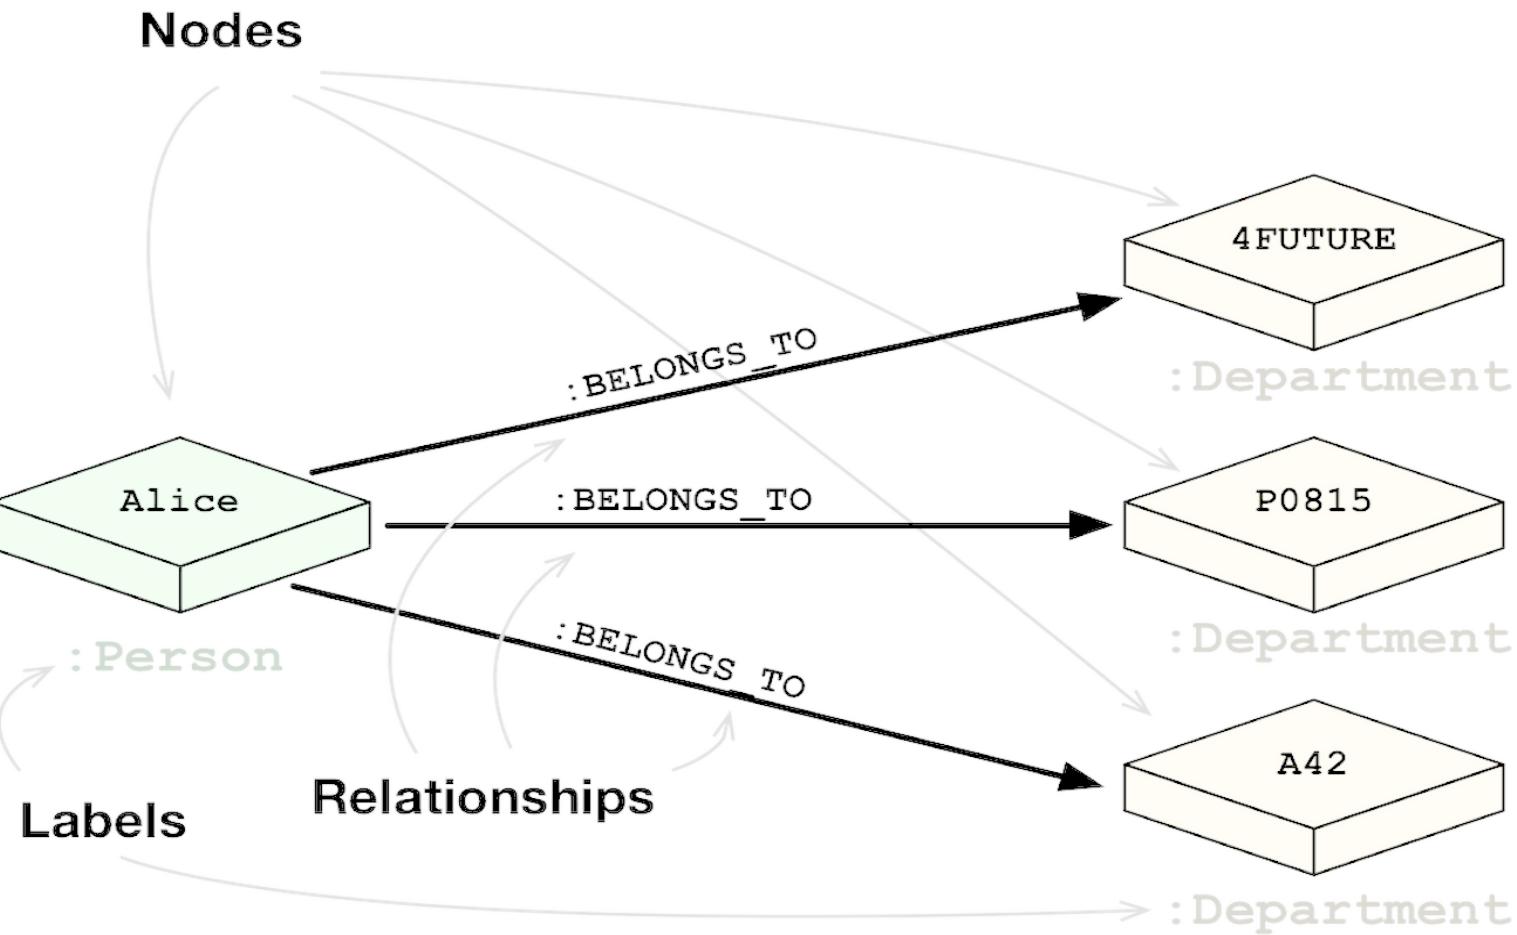
Graphs & Relational Databases

Relational Databases

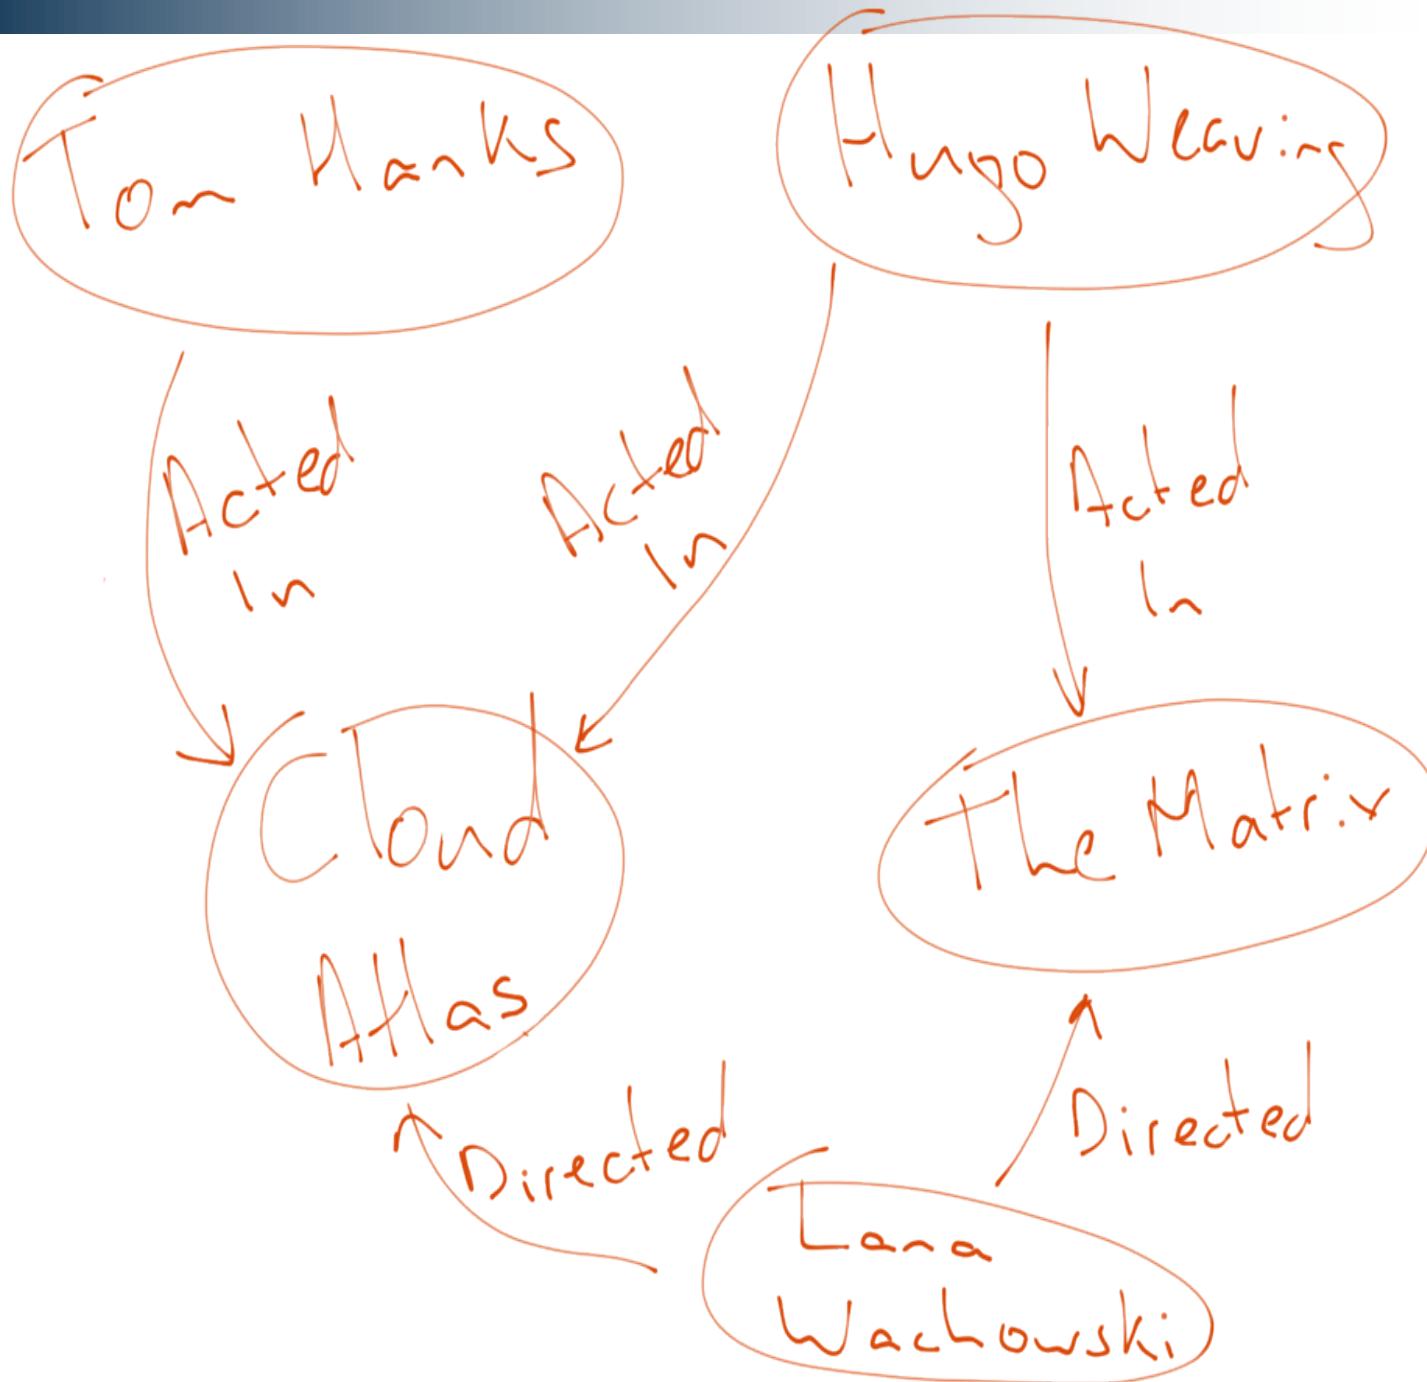


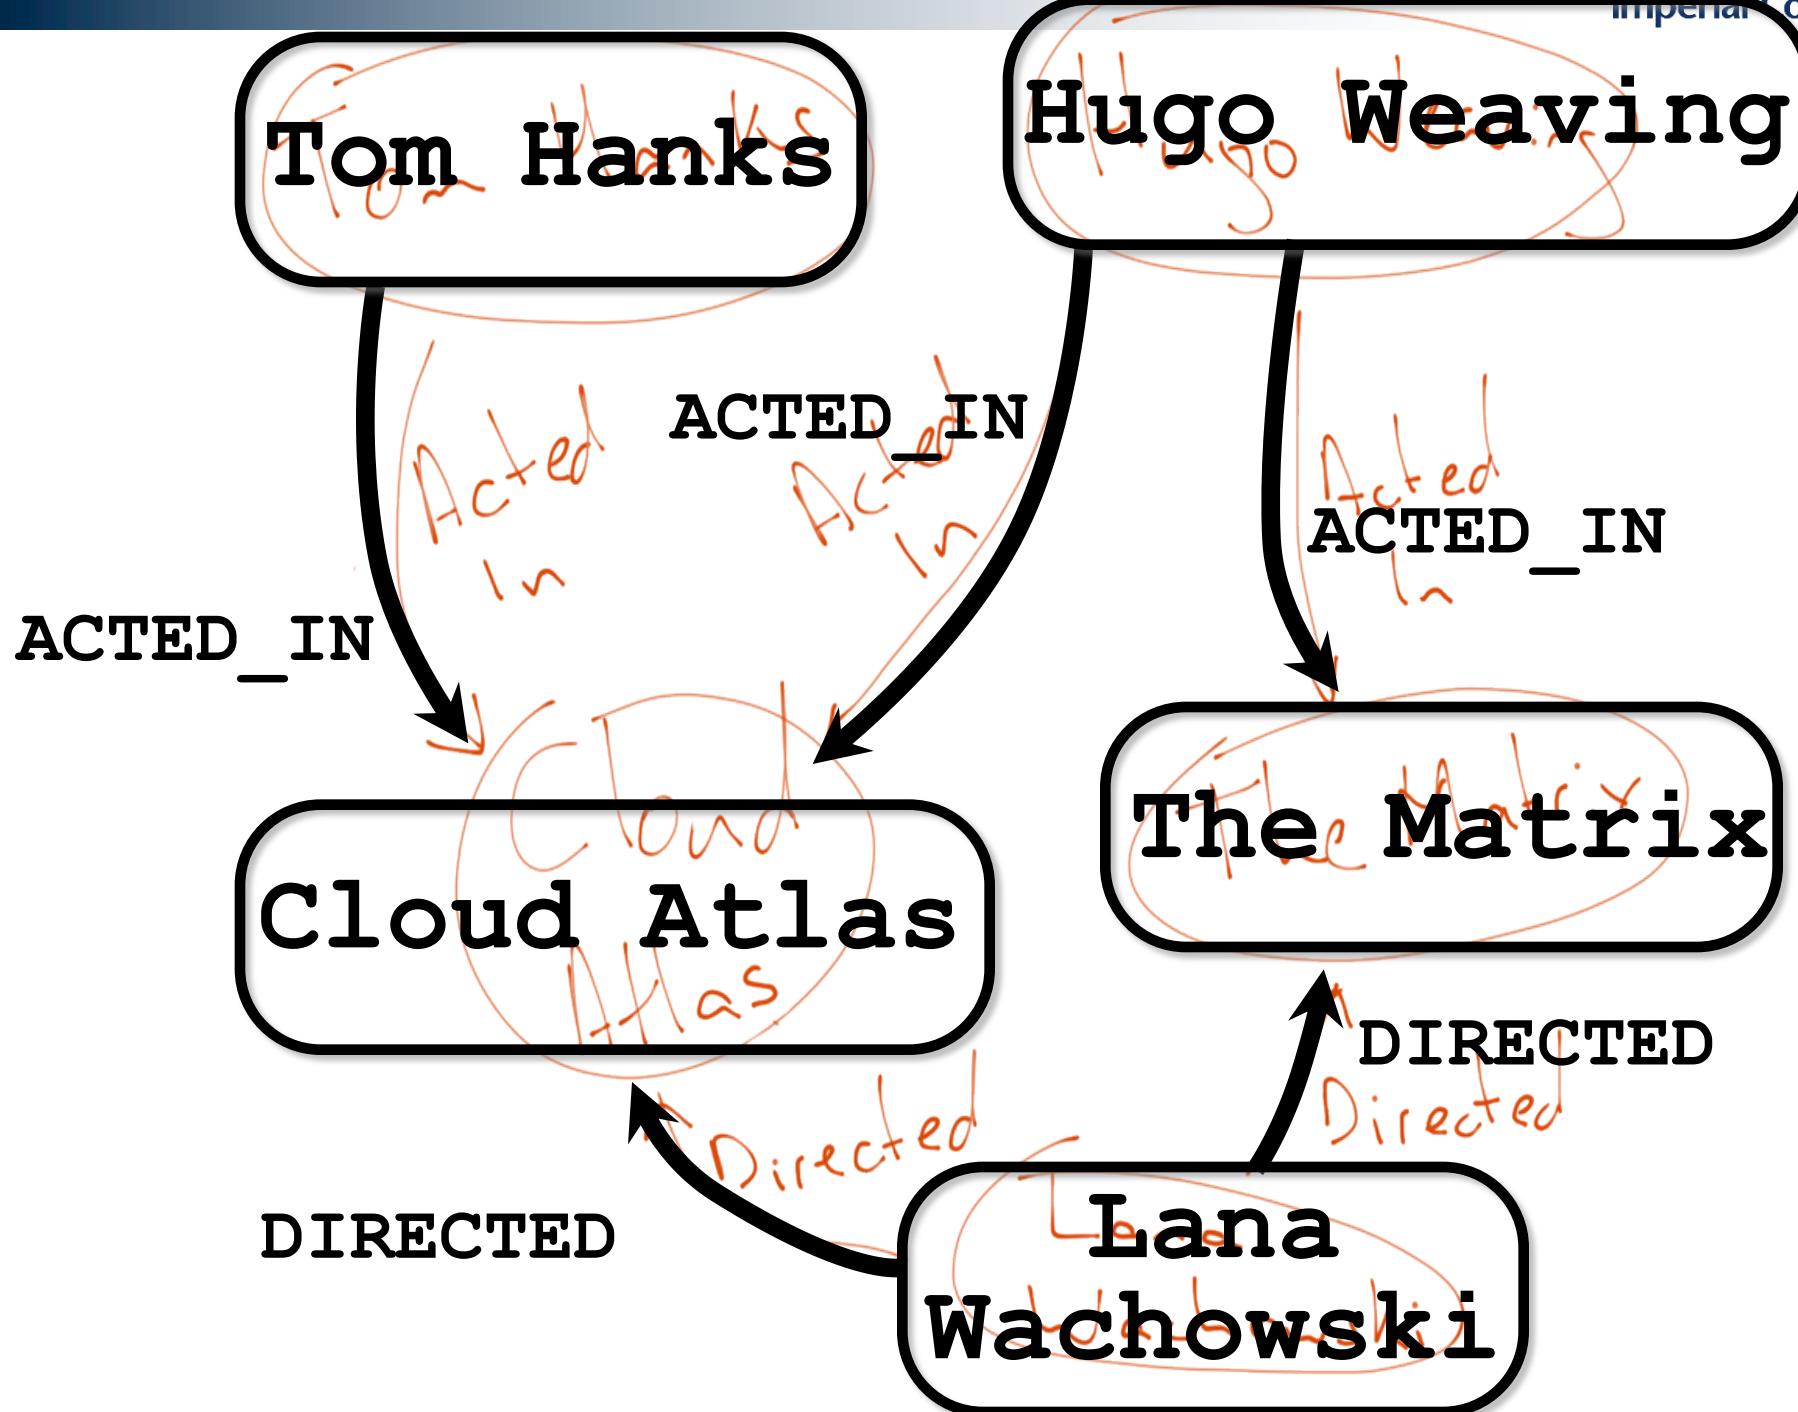
Graph Databases

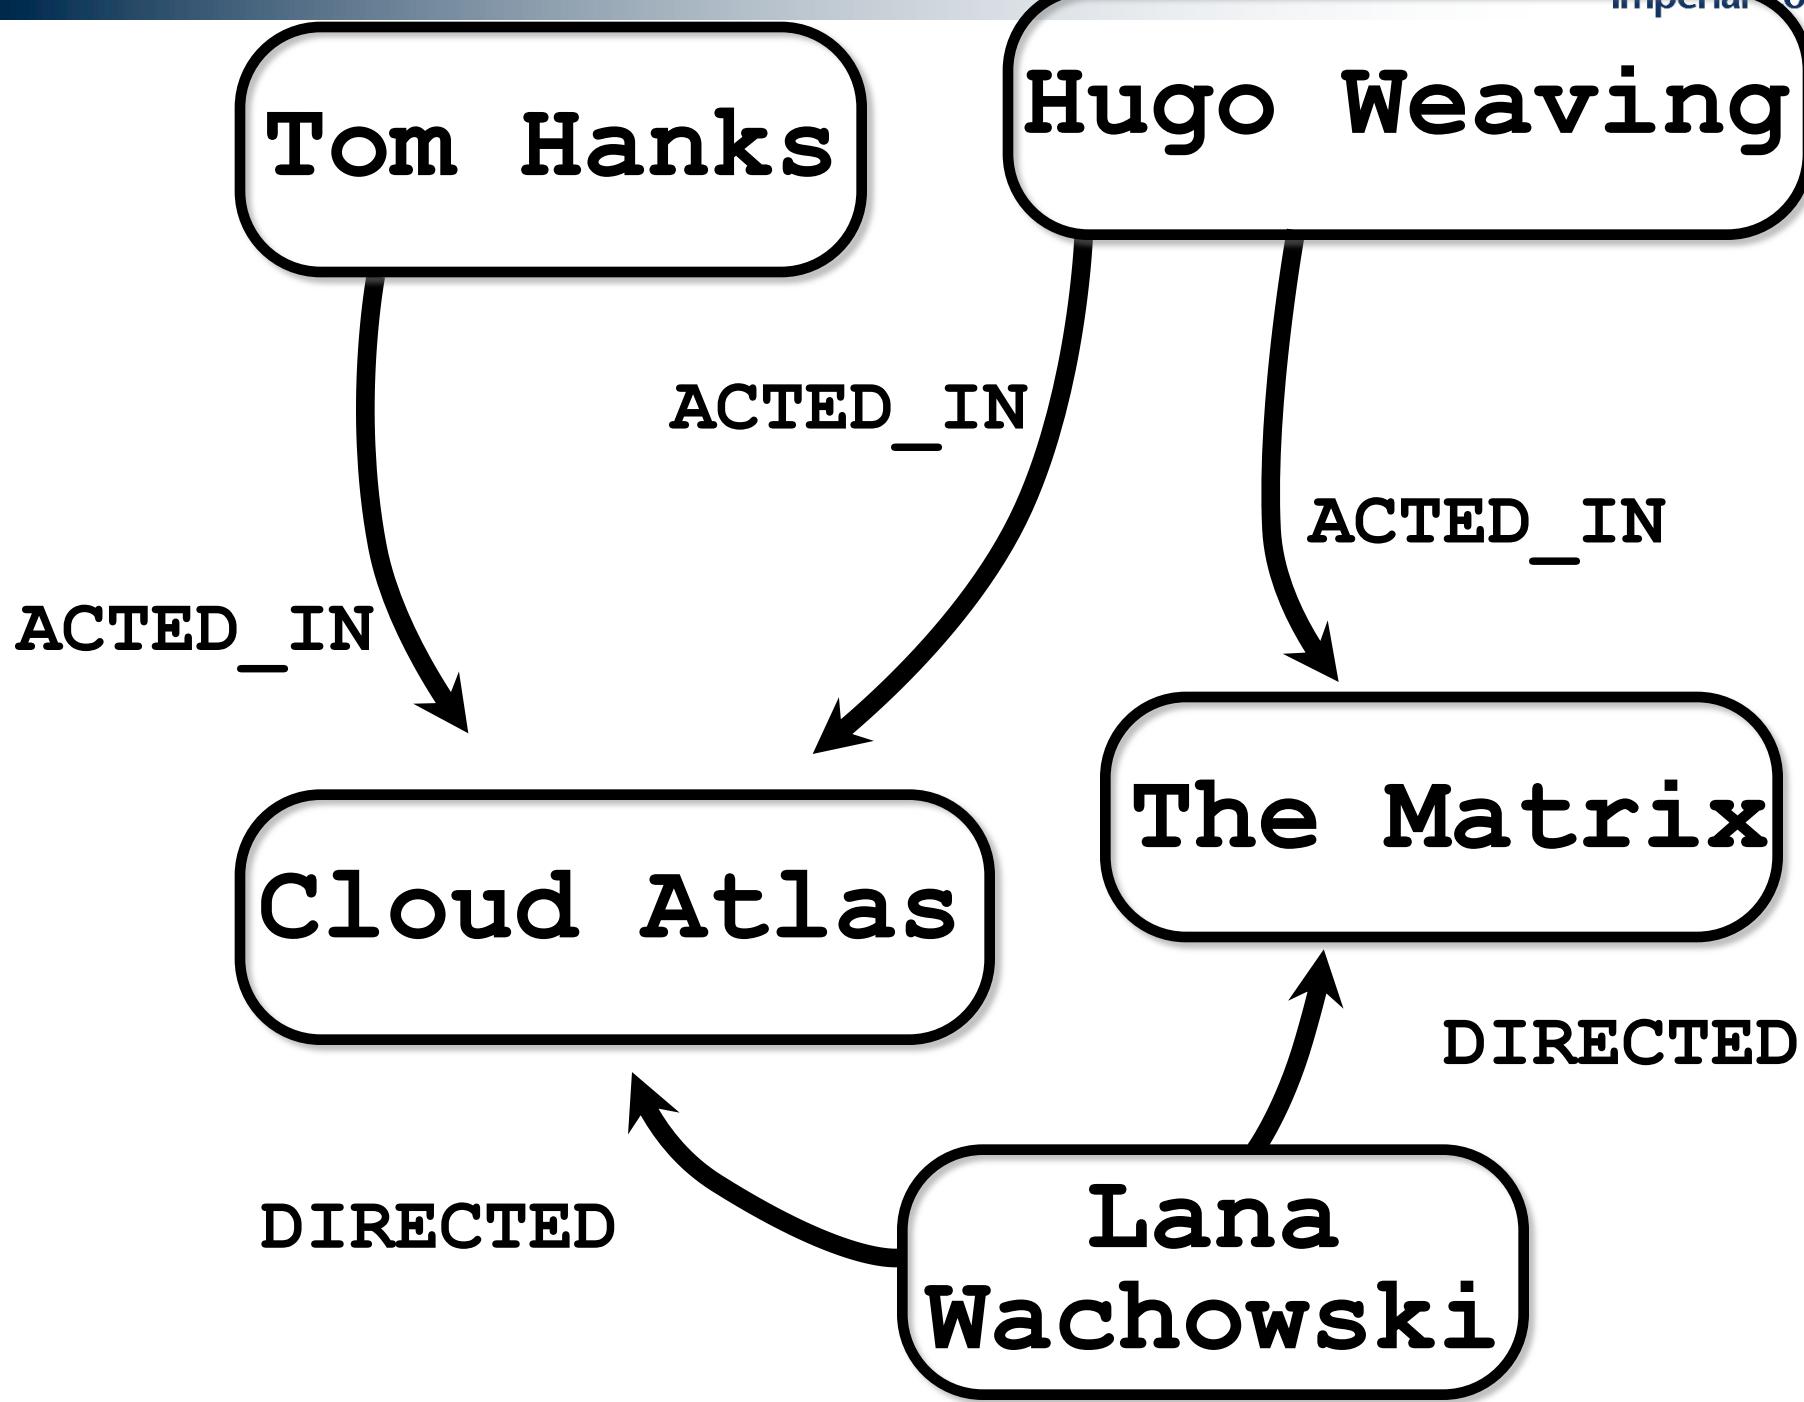




Modeling with Graphs







name: Tom Hanks
nationality: USA
won: Oscar, Emmy

name: Hugo Weaving
nationality: Australia
won: MTV Movie Award

ACTED_IN
role: Bill Smoke

ACTED_IN
role: Agent Smith

ACTED_IN
role: Zachry

title: Cloud Atlas
genre: drama, sci-fi

title: The Matrix
genre: sci-fi

DIRECTED

DIRECTED

name: Lana Wachowski
nationality: USA
won: Razzie, Hugo

Overview of Neo4j

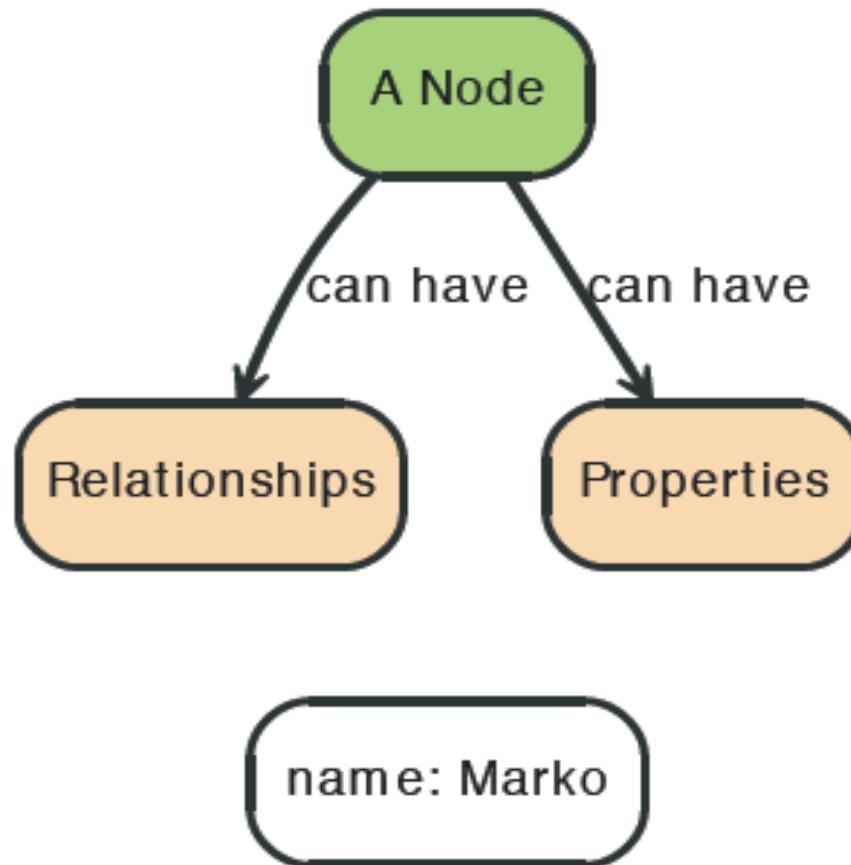
What is a Graph Database?

- A database with an explicit graph structure
- Each node knows its adjacent nodes
- As the number of nodes increases, the cost of a local step (or hop) remains the same
- Plus an Index for lookups

Translating to Neo4j

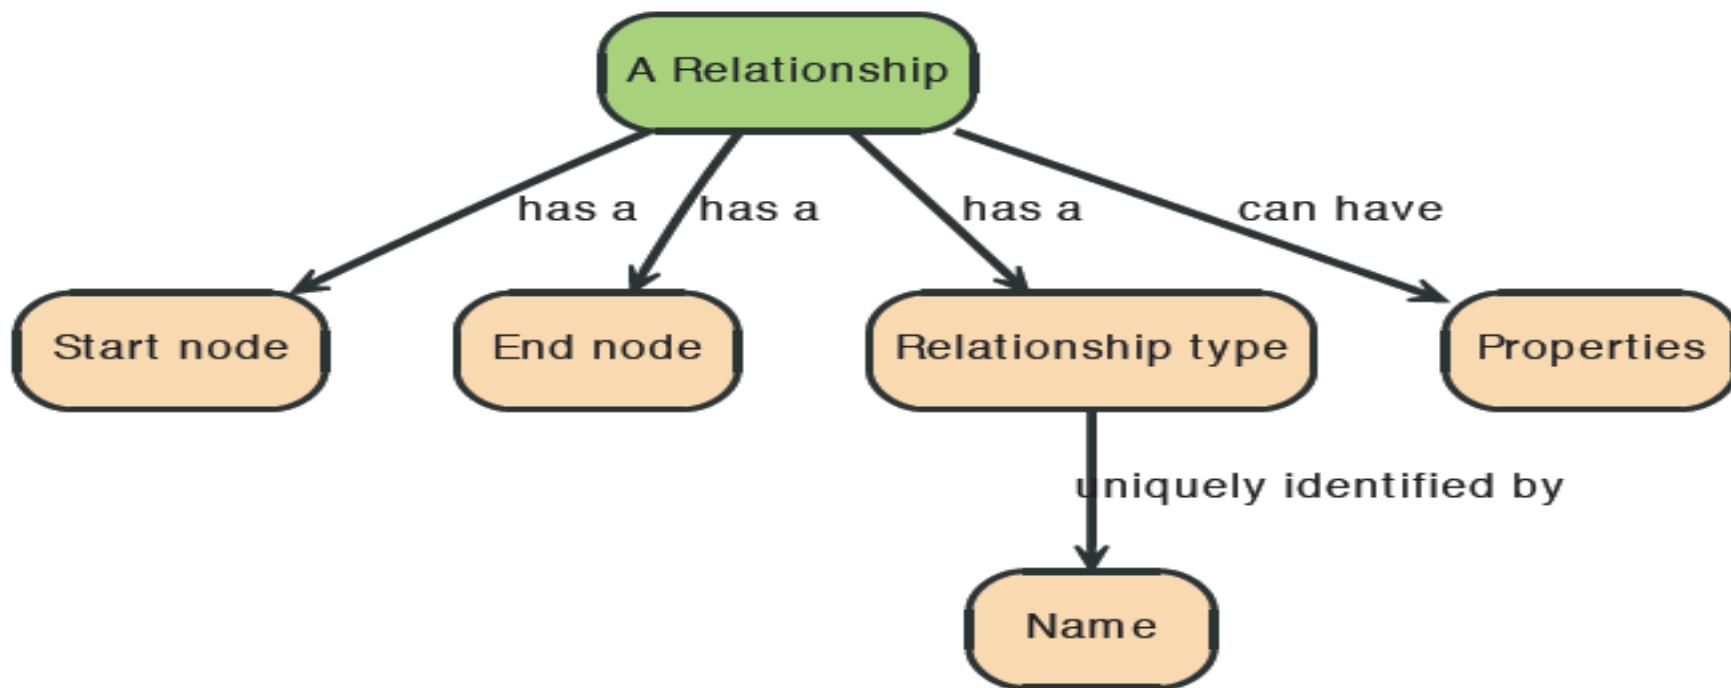
- Each entity table is represented by a label on nodes
- Each row in a entity table is a node
- Columns on those tables become node properties.
- Add unique constraints for business primary keys, add indexes for frequent lookup attributes

Node in Neo4j

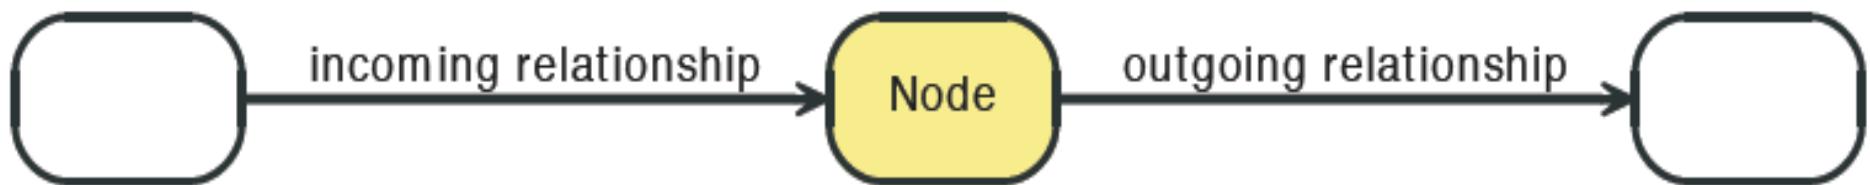
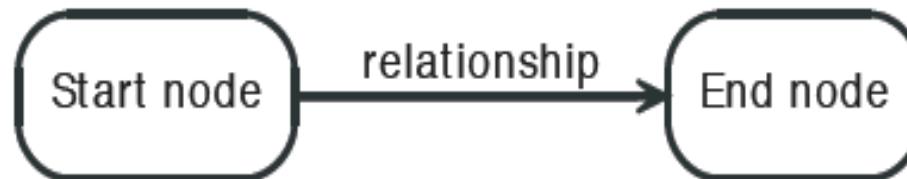


Relationships in Neo4j

Relationships between nodes are a key part of Neo4j.



Relationships in Neo4j

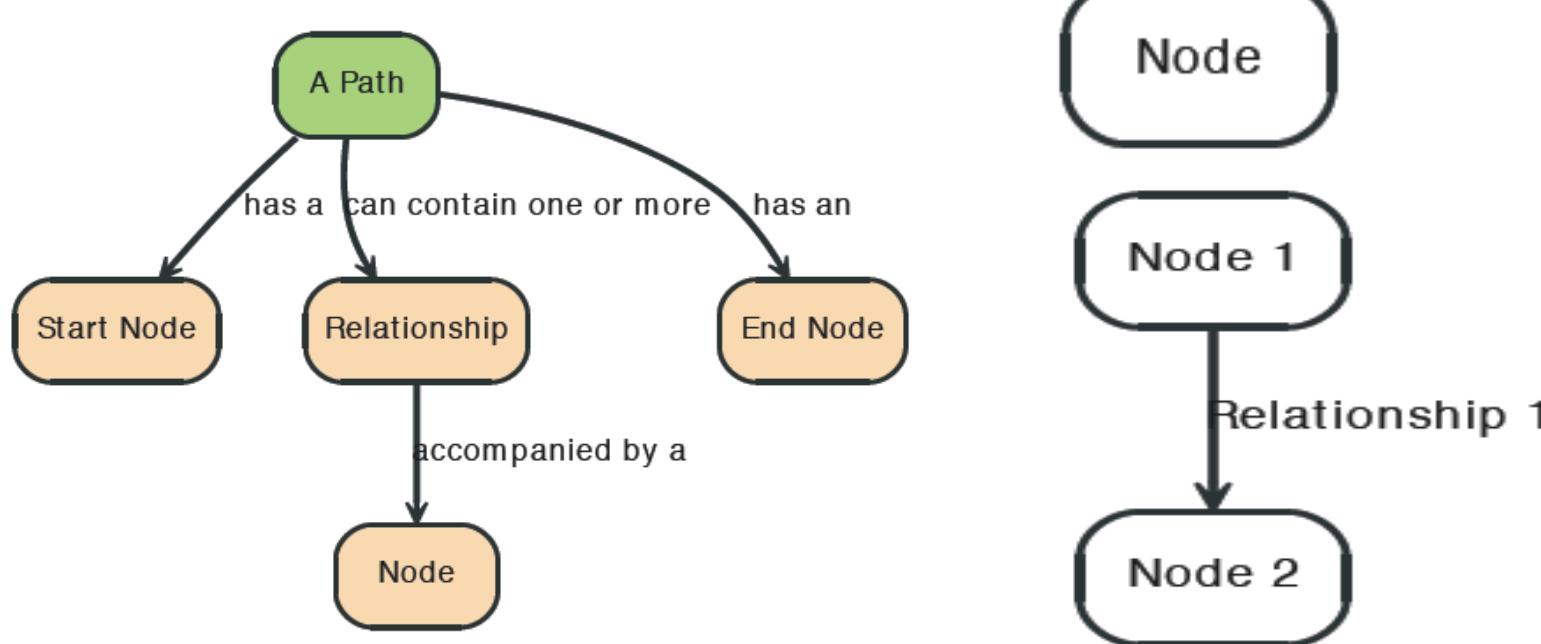


Properties

- Both nodes and relationships can have properties.
- Properties are key-value pairs where the key is a string.
- Property values can be either a primitive or an array of one primitive type.
- For example String, int and int[] values are valid for properties.

Paths in Neo4j

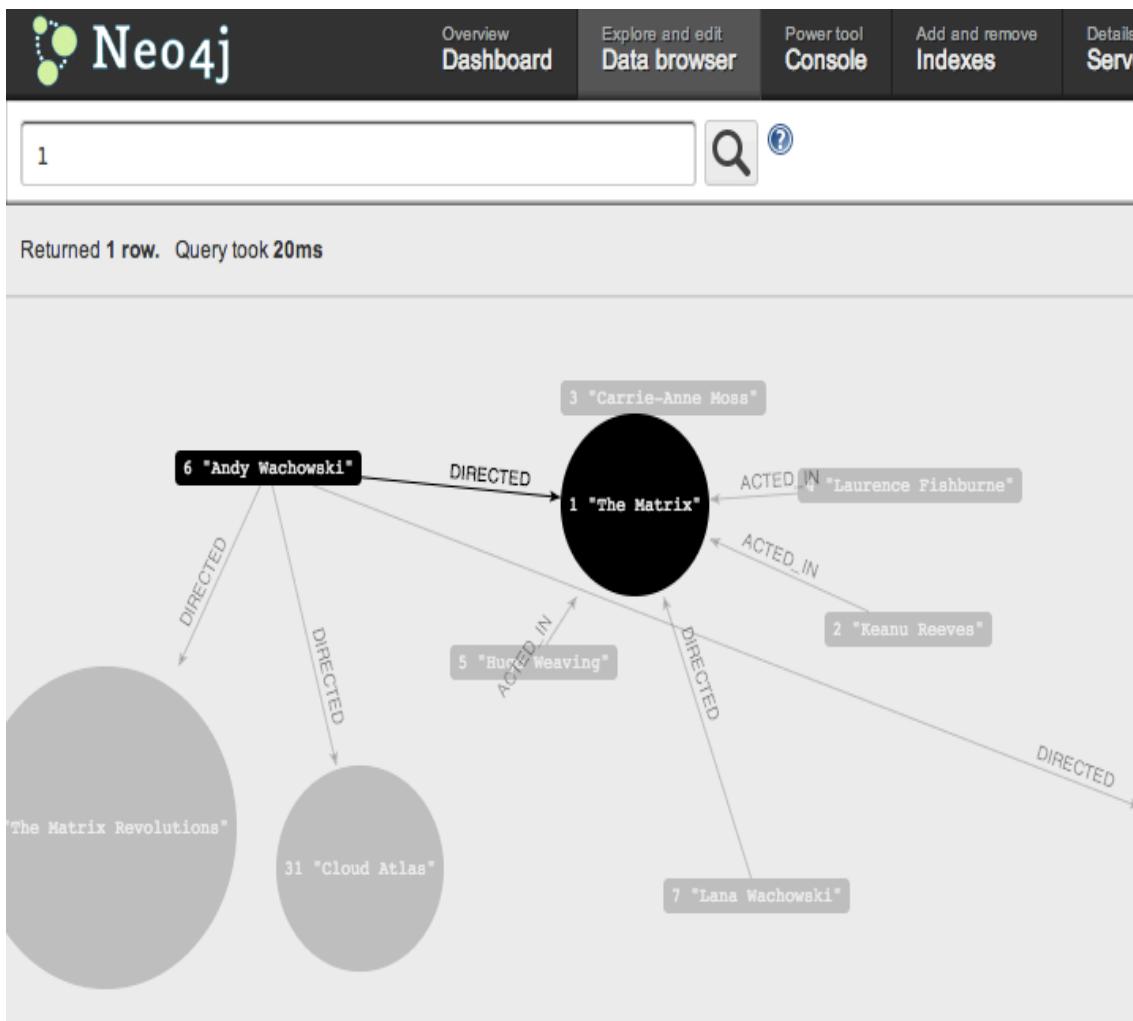
A path is one or more nodes with connecting relationships, typically retrieved as a query or traversal result.



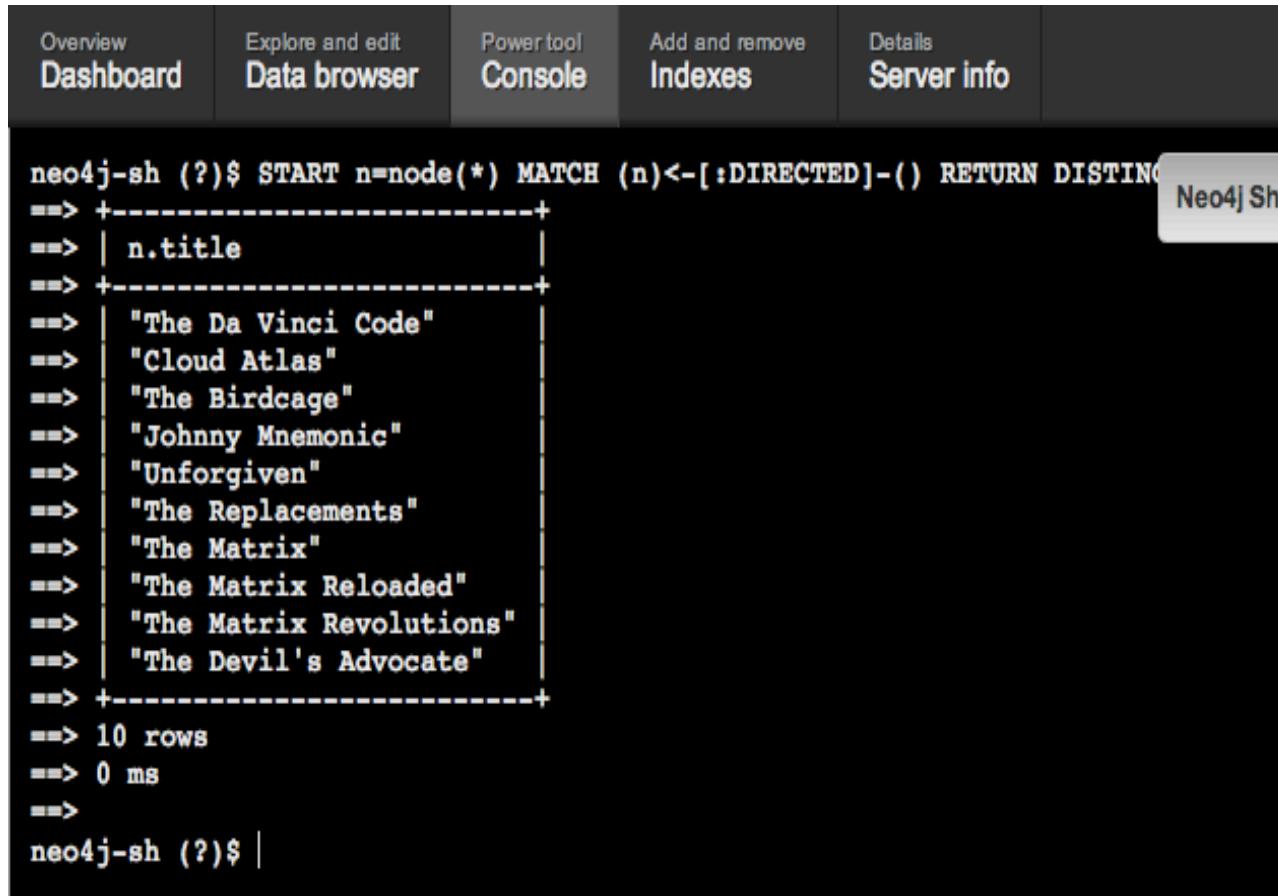
Test Server

<https://neo4j.com/lp/try-neo4j-sandbox/>

Neo4j Data Browser



Neo4j Console



The screenshot shows the Neo4j Console interface with a dark theme. The top navigation bar includes links for Overview, Dashboard, Explore and edit Data browser, Power tool Console (which is currently selected), Add and remove Indexes, Details, and Server info. The main area displays a command-line session:

```
neo4j-sh (?)$ START n=node(*) MATCH (n)<-[ :DIRECTED ] -() RETURN DISTINCT n.title
==> +-----+
==> | n.title
==> +-----+
==> | "The Da Vinci Code"
==> | "Cloud Atlas"
==> | "The Birdcage"
==> | "Johnny Mnemonic"
==> | "Unforgiven"
==> | "The Replacements"
==> | "The Matrix"
==> | "The Matrix Reloaded"
==> | "The Matrix Revolutions"
==> | "The Devil's Advocate"
==> +-----+
==> 10 rows
==> 0 ms
==>
neo4j-sh (?)$ |
```

A tooltip labeled "Neo4j Shell" is visible near the right edge of the screen.

Introduction to Cypher

- Declarative Pattern-Matching language
- SQL-like syntax
- Designed for graphs

Two nodes, one relationship

```
START a=node(*)
```

```
MATCH (a) - ->(b)
```

```
RETURN a, b;
```



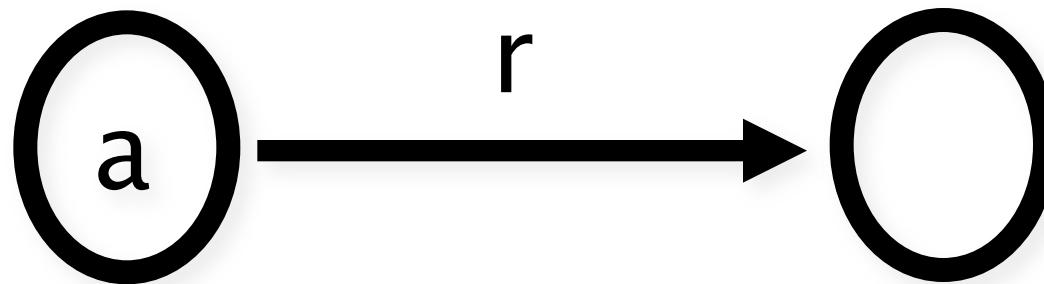
Two nodes, one relationship

```
START a=node(*)  
MATCH (a) - ->()  
RETURN a;
```



```
START a=node(*)  
MATCH (a) - ->()  
RETURN a.name;
```

Two nodes, one relationship

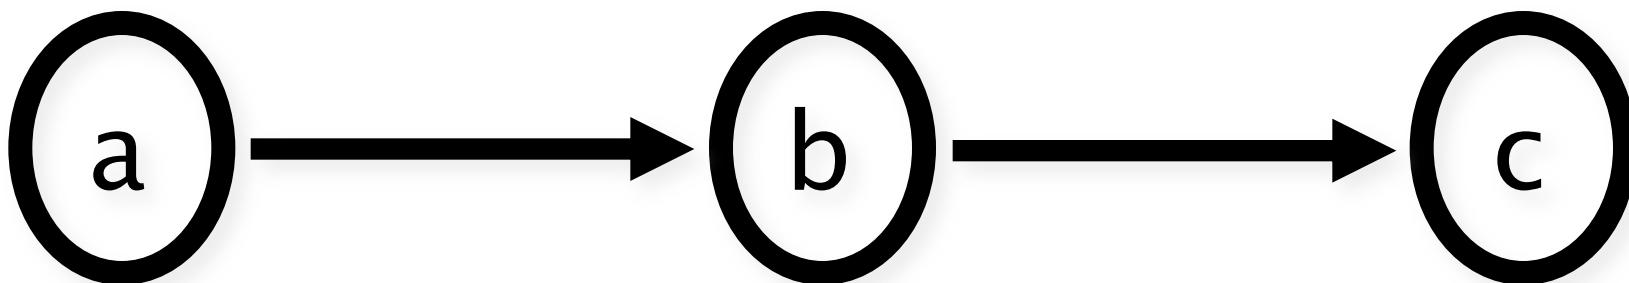


(a) -[r]-> ()



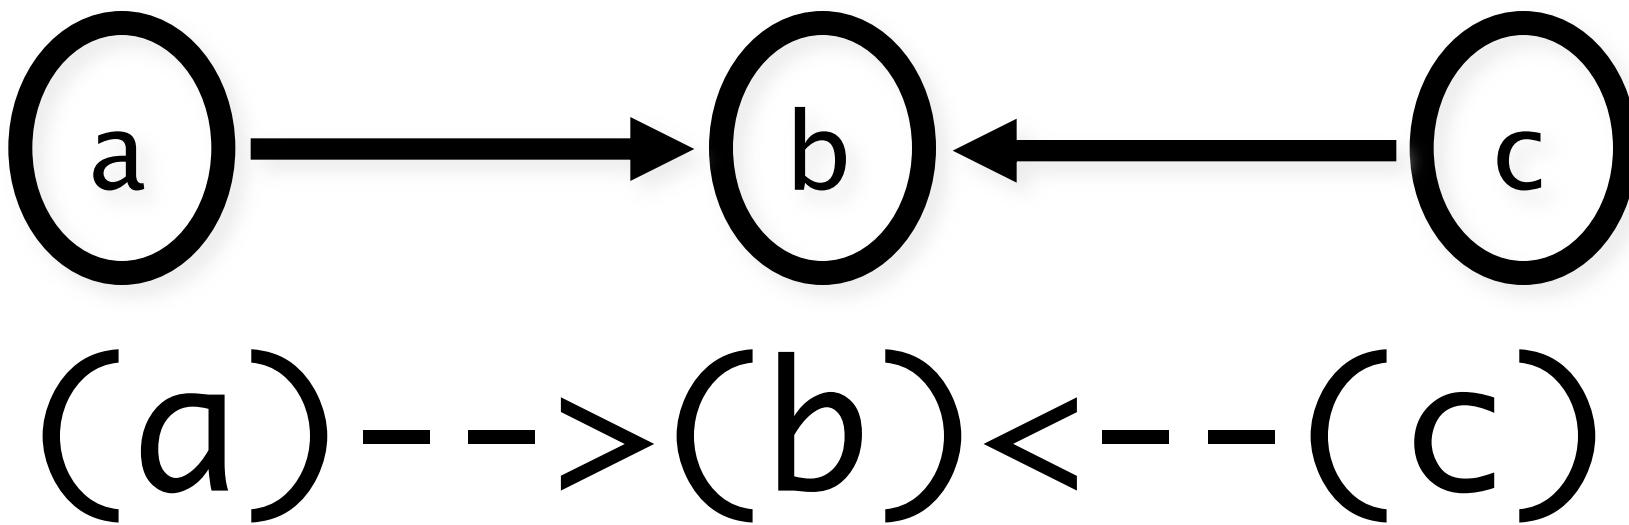
(a) -[:ACTED_IN]-> (m)

Paths

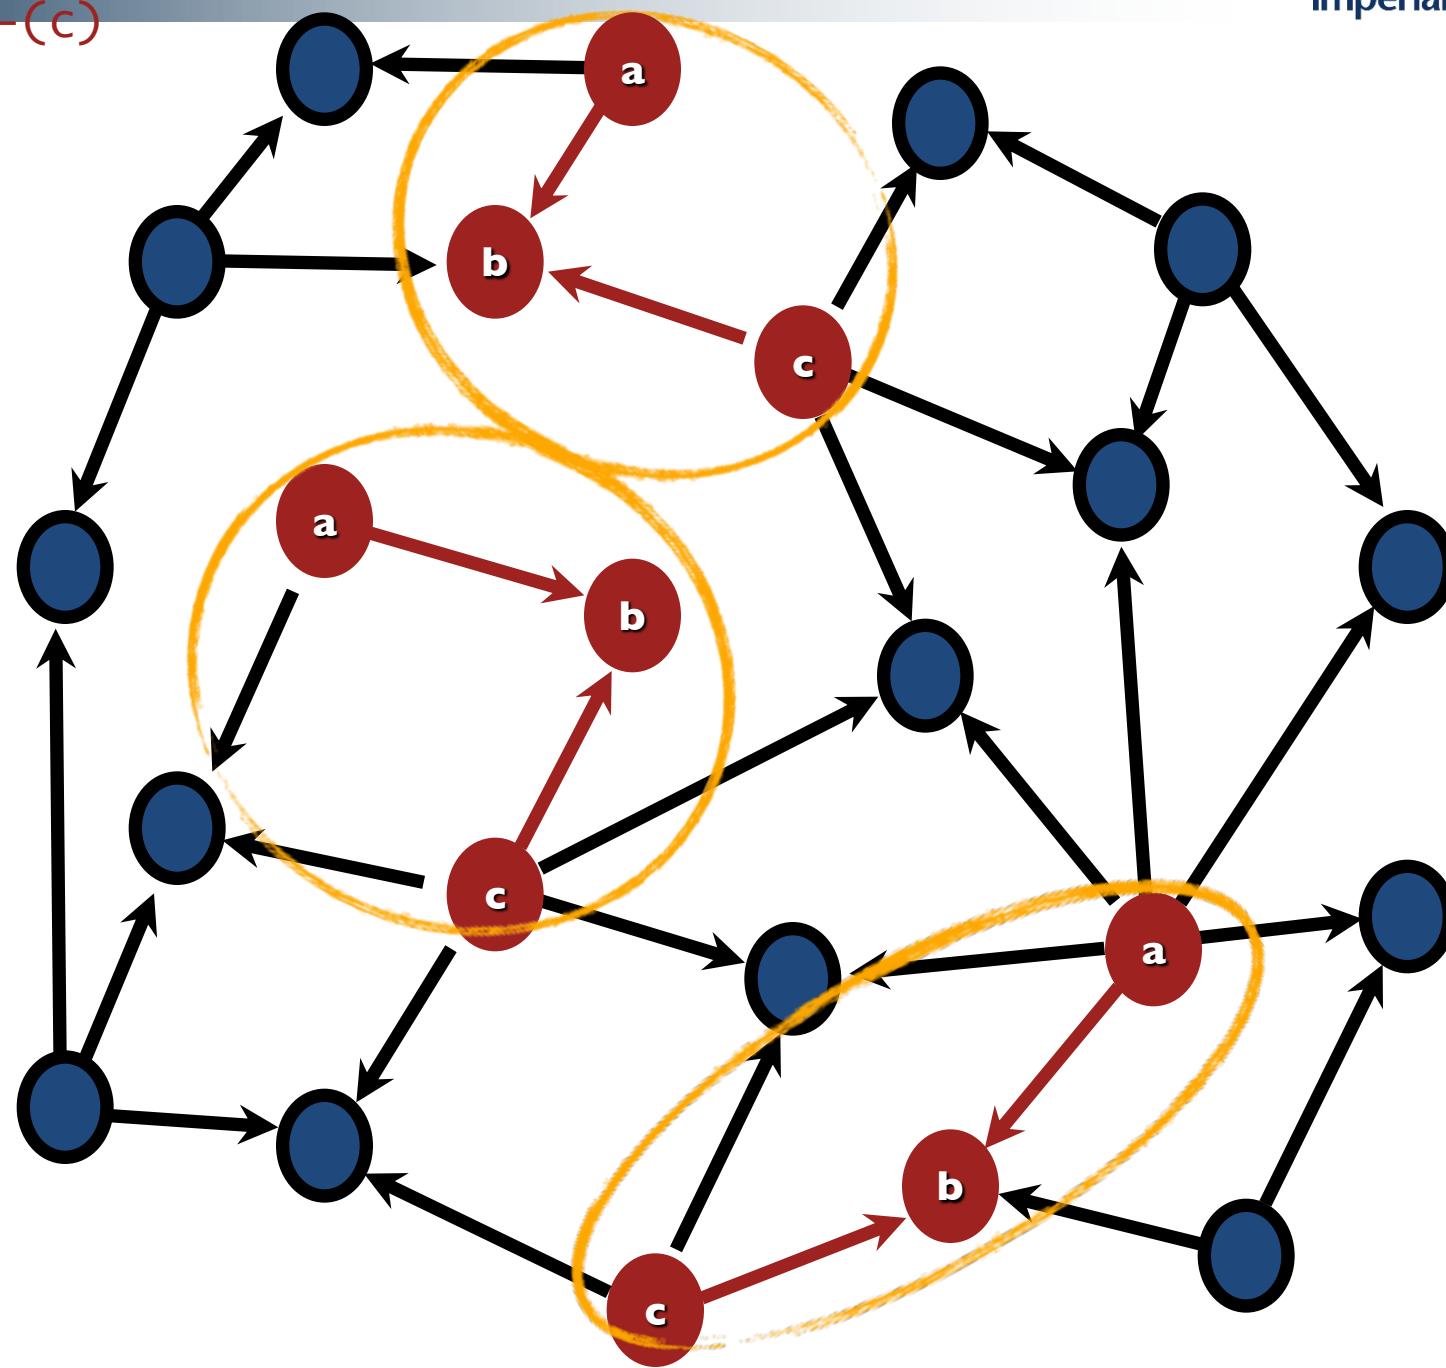


(a)--->(b)--->(c)

Paths



(a)--->(b)<--(c)



Paths

```
START a=node(*)  
MATCH (a) - [:ACTED_IN] ->(m)< - [:DIRECTED] - (d)  
RETURN a.name, m.title, d.name;
```

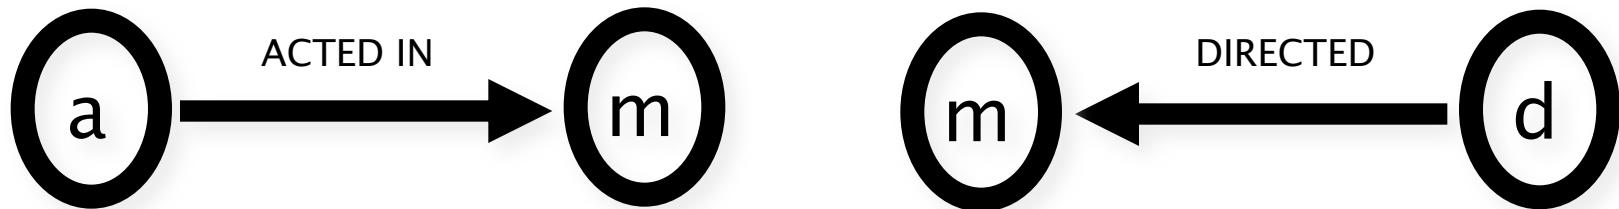
a.name	m.title	d.name
“Keanu Reeves”	“The Matrix”	“Andy Wachowski”
“Keanu Reeves”	“The Matrix Reloaded”	“Andy Wachowski”
“Noah Wyle”	“A Few Good Men”	“Rob Reiner”
“Tom Hanks”	“Cloud Atlas”	“Andy Wachowski”
...

Paths

```
START a=node(*)
```

```
MATCH (a)-[:ACTED_IN]->(m), (m)<-[DIRECTED]-(d)
```

```
RETURN a.name, m.title, d.name;
```



```
START a=node(*)
```

```
MATCH (a)-[:ACTED_IN]->(m),
```

```
(d)-[:DIRECTED]->(m)
```

```
RETURN a.name, m.title, d.name;
```

Sort & Limit

```
START a=node(*)MATCH (a)-[:ACTED_IN]->(m)<-
[:DIRECTED]-(d)RETURN a.name, d.name,
count(*) AS count
ORDER BY(count) DESC
LIMIT 5;
```

Aggregation

count(x) - add up the number of occurrences

min(x) - get the lowest value

max(x) - get the highest value

avg(x) - get the average of a numeric value

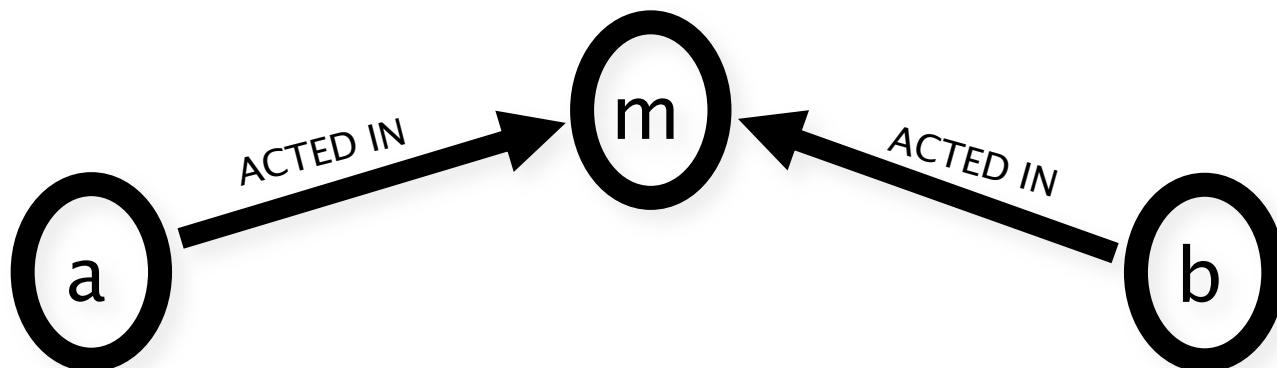
collect(x) - collected all the occurrences into an array

Aggregation

```
START a=node(*)
```

```
MATCH (a)-[:ACTED_IN]->(m)<-[:DIRECTED]-(d)
```

```
RETURN a.name, d.name, collect(m.title);
```



Find a specific node (all-node query)

```
START n=node(*)  
WHERE has(n.name) AND n.name = "Tom Hanks"  
RETURN n;
```

WHERE - filter the results

has(n.name) - the name property must exist

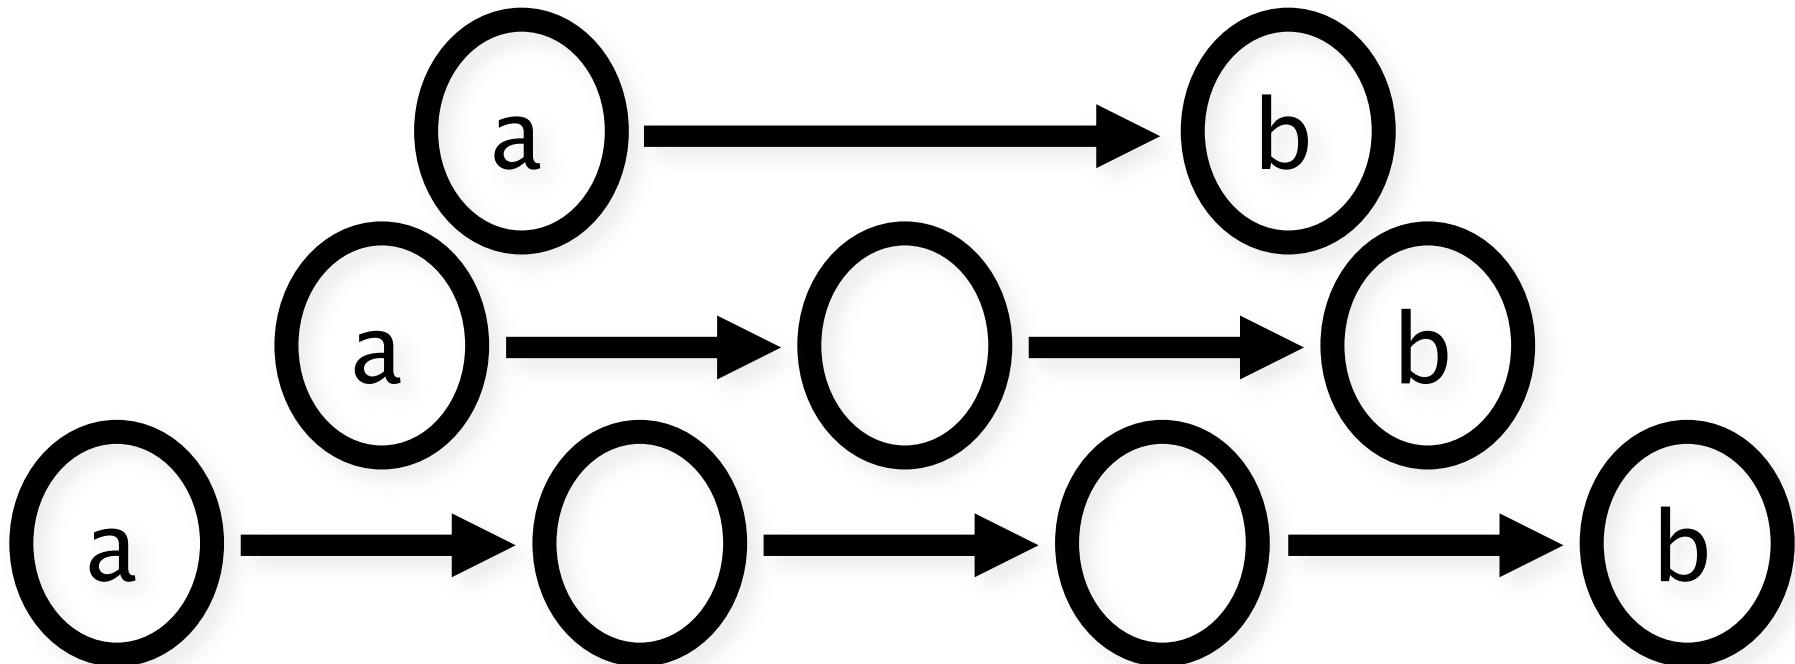
n.name = "Tom Hanks" - and have that value

Matching multiple relationships

```
START a=node(*)  
MATCH (a) - [:ACTED_IN|DIRECTED] -> () <-  
[:ACTED_IN|DIRECTED] - (b)  
CREATE UNIQUE (a) - [:KNOWS] -> (b);
```

(Create KNOWS relationships between anyone, Actors or Directors, who worked together)

Variable length paths



(a)-[*1..3]->(b)

Friends-of-Friends

```
MATCH (keanu) - [:KNOWS*2] ->(fof)
WHERE keanu.name = "Keanu Reeves"
RETURN DISTINCT fof.name;
```

Conditions

Constraints on properties

```
MATCH (actor)-[r:ACTED_IN]->(movie)
WHERE "Neo" IN r.roles AND actor.name = "Keanu
Reeves"
RETURN DISTINCT movie.title;
```

(Movies in which Keanu Reeves played Neo)

Constraints based on patterns

```
MATCH (gene) - [:ACTED_IN] -> (movie) <- [:ACTED_IN] - (n)
WHERE gene.name = "Gene Hackman"
RETURN DISTINCT n.name;
```

(Actors who worked with Gene Hackman)

Constraints based on patterns

```
MATCH (gene) - [:ACTED_IN] -> (movie) <- [:ACTED_IN] - (n)  
WHERE (n) - [:DIRECTED] -> () AND gene.name = "Gene  
Hackman"
```

```
RETURN DISTINCT n.name;
```

(Actors who worked with Gene and were directors)

Updating Graphs with Cypher

Creating nodes

```
CREATE ({title:"Mystic River", released:1993});
```

```
START a=node(*)  
MATCH (a)  
WHERE a.title = "Mystic River"  
RETURN a
```

Adding properties

```
START a=node(*)  
MATCH (a)  
WHERE a.title = "Mystic River"  
SET movie.tagline = "We bury our sins here, Dave. We wash them  
clean."  
RETURN movie;
```

Creating Relationships

```
CREATE UNIQUE (kevin)-[:ACTED_IN {roles:["Sean"]}]>(movie)  
WHERE movie.title = "Mystic River" AND kevin.name = "Kevin Bacon"
```

```
MATCH (kevin)-[:ACTED_IN]->(movie)  
WHERE kevin.name = "Kevin Bacon"  
RETURN DISTINCT movie.title;
```

Deleting relationships

```
MATCH (emil)-[r]->()
WHERE emil.name = "Emil Eifrem"
DELETE r;
```

Deleting nodes and relationships

```
MATCH (emil)-[r]->()
WHERE emil.name = "Emil Eifrem"
DELETE r, emil;
```

Deleting nodes and relationships

```
MATCH (emil)-[r?]->()
WHERE emil.name = "Emil Eifrem"
DELETE r, emil;
```